

PCSE OpenMP Homework 1, part 2 (hw1_part2)

(updated 3/7/2017 – see green text, pg 4 – Fortraner use f90_setaffinity)

Overview

This is part2 of 2 parts for hw1. You will parallelize an existing the rb.c or rb.F90 serial program using OMP directives and report scaling information.

Work with either C or Fortran codes—NOT both. DO THIS ASSIGNMENT on Stampede.

Resources:

Untar the file omp_hw1_part2.tar into a directory with the following commands:

```
cd
tar xvf ~train00/omp_hw1_part2.tar
cd hw1/rb
```

General:

You can use the login node testing and parallelizing codes, but use **idev** to interactively run an executable when benchmarking. You can check where the threads are running by checking the load on each core with the **top** utility.

top # then hit the 1 character

The basic steps are.

1. Look over the serial code.
2. Copy the serial code to another file (with a specific name) and parallelize it using OpenMP directives. Compile the code
3. Run the code
4. **Plot and analyze results, as requested.**

To make this a bit easier, scripts have been made for you.
For compiling execute:

```
./compile my_prog.c or ./compile my_prog.F90
```

The **compile** script will make an executable named **a.out_my_prog**.

To interactively run an executable with various number of threads, execute:

```
./dothis # (will run parallel code for different # of threads)
```

FYI: The timer.c and affinity.c codes are compiled with the “-c” option in the **compile** script and creates just “object” files (timer.o, affinity.o). These are “loaded” into your code on the compile line like this:

```
ifort -qopenmp timer.o affinity.o prb_d.F90 -o a.out_prb_d  
icc -qopenmp timer.o affinity.o prb_d.c -o a.out_prb_d
```

(You can always type these on the command line instead of using **compile**.)

Instructions hw1 – part 2

1. The red-black (rb) program is an example of an update mechanism that might be encountered in various algorithms. In a red-black algorithm, red or black elements of an array are updated separately. Two red-black loops update array "a" with other elements of "a". In this red-black algorithm the "red" elements of the "a" array, {1,3,5,7,...}, are updated, and then the "black" elements of the array, {0,2,4,...}, are updated. (Here, we could have simply called them odd and even, respectively. (But, of course, in 2-D or 3-D algorithms, associating color with a pattern makes much more sense for visualization.)

The objective here is to first parallelize loops in a simple manner (part a), and then reduce the number of parallel forks in part b. In part c you will also try to optimized the code and check the performance of binding (affinity) schemes.

You will create 3 different parallel codes from the same serial code source, parallelize them in slightly different ways, and then benchmark and analyze them. In the [Report show the OpenMP changes to the pseudo code \(below\), and include any loop changes in the optimization section \(part c\)](#). Create required plots, and discuss the behavior and/or parallelizability where specified. Report times for the while loop only, **as shown in the serial code, rb.c**.

```
Begin Program   Pseudo Code
<any new variables here>

#pragma omp parallel ... or !$omp parallel
  nt = omp_get_num_threads(); WARMUP,initial parallel

#pragma omp ... or !$omp ...
  Initialization:

  Start timer
  while[

    #pragma omp ... or !$omp ...
    1st red-black loop (if you change

    #pragma omp ... or !$omp ...
    2st red-black loop

    #pragma omp ... or !$omp ...
    error var initialization

    #pragma omp ... or !$omp ...
    error loop

  ]
  End timer

  Report time

END
```

Use the **schedule(runtime)** clause on the work sharing loops in the while loop so that they can run with different scheduling algorithms (static and dynamic).

Use the **compile** and **dothis** scripts for compiling code and executing them. Execute the **dothis** script interactively on a compute node – use idev.

- a. Parallelize the code in the most obvious/simple (and correct) way—by separately parallelizing each loop (with **parallel do** or **parallel for**) within the timed region. General Hint: be careful with initializing the error var).

- i. Copy rb.c/F90 to prb_a.c or prb_a.F90
- ii. Parallelize all loops in with “parallel for” or “parallel do”, EXCEPT for the initialization.
- iii. Compile: **./compile prb_a.c** (creates a.out_prb_a)
- iv. Execute **./dothis** to execute a.out_prb_a (various # of threads and schedules). Times for 1-8 threads are reported for 4 types of scheduling.
Report: Do the various Scheduling methods make much difference? How good/bad is the scaling up to 8 threads?
- v. Change code to now parallelize the initialization loop. DO NOT include a schedule clause—let it default to static scheduling. Compile with **compile prb_a.c** and execute **dothis**.
Report: How much better the scaling is? Why is it better? Do the various Scheduling methods now make much difference?
- vi. Watch where the threads are executing with top. (On the node in another window execute “top” and then hit the “1” key to see the loads. Type s and then 1 to have top update every second.)
Report: Is there any consistency to where the threads are running in iv & v?

Report: Show coding changes. Plot the scaling for the STATIC case (without a chunk size) with and without the parallelization of the initialization loop (label as ser_init and par_init in plot).

- b. Using prc_a.c, reduce the number of parallel regions (directives).

- i. copy prb_a.c/F90 to prb_b.c or prb_b.F90
- ii. Modify directives in the while loop so that there is only one parallel directive in the while loop. (Hint: Think “single” for the error and niter vars.)
- iii. Compile: use **compile prb_b.c**; and execute: **dothis**.

Report: Show coding changes. Does the scheduling make any difference? Does the code scale better or worse than the runs in part a; does it run faster? Give a reason why they “should” be about the same, even though the number of forking requests has been significantly reduced.

- c. There are 2 optimizations to be performed here: 1.) Use a single parallel region (directive) outside the while loop; 2.) If the 2 red-black loops can be combined, fuse them. We use 16 threads here.

General Instructions: Copy `rb.c/F90` to `prb_c.c` or `prb_c.F90` [or copy from `prb_b.c/F90`], parallelize, compile with ***compile prb_c.c/F90***; use ***dothis*** to execute and run on a compute node.

- i. Once you have constructed a **single parallel directive outside** of the while loop, compile and run multiple times with ***dothis***, save output. **(Include the initialization of the `a` array in the parallel region.)**
- ii. Now combine red-black loops, compile and run multiple times with ***dothis*** (save output).

Review the contents of ***dothis*** and ***compile*** scripts.