

Home work instructions

Stampede2 --- Skylake Queues

I forgot to explicitly stating this in the documents that I posted yesterday. I apologize for the confusion.

As you may remember, Stampede2 has compute nodes with 2 architectures: Knights Landing (KNL) nodes and **Skylake nodes**. Wherever it says 'compute nodes' in the documents please assume '**Skylake compute nodes**'. All the output, tests, and timings have to be produced on the Skylake nodes.

Please understand that the default queues are referring to the KNL nodes, i.e. to the nodes that you should not be using. Therefore **you have to specify a queue** when you are submitting a job or starting an interactive session.

Interactive session with idev:

```
idev -p skx-dev
```

```
idev -p skx-normal
```

In a job script add one of these lines

```
#SBATCH -p skx-dev
```

```
#SBATCH -p skx-normal
```

To check the queue names:

```
sinfo -n asdfasdf
```

Other commands (just a reminder)

```
showq -l
```

```
showq -l -u
```

```
scancel <job-id>
```

Note that instructions are only for Stampede2 and are not relevant if you are using Maverick (which you should only do in an emergency)

Best, Lars