

Programming with MPI

Basic send and receive

Jan Thorbecke

Acknowledgments

- This course is partly based on the MPI course developed by
 - Rolf Rabenseifner at the High-Performance Computing-Center Stuttgart ([HLRS](http://www.hlrs.de/home/)), University of Stuttgart in collaboration with the EPCC Training and Education Centre, Edinburgh Parallel Computing Centre, University of Edinburgh.
<http://www.hlrs.de/home/>
- [CSC](https://www.csc.fi) – IT Center for Science Ltd.
<https://www.csc.fi>



CSC-IT CENTER FOR SCIENCE



High-Performance Computing Center | Stuttgart

Contents

- Initialisation of MPI
 - exercise: HelloWorld
- Basic Send & Recv
 - exercise: Sum
 - exercise: SendRecv
- more Send Receive messages
 - exercise: PingPong (optional)
 - exercise: Ring (optional)

A Minimal MPI Program (C)

```
#include "mpi.h"
#include <stdio.h>

int main( int argc, char *argv[] )
{
    err = MPI_Init( &argc, &argv );
    printf( "Hello, world!\n" );
    MPI_Finalize();
    return 0;
}
```

A Minimal MPI Program (Fortran 90)

```
program main
use MPI
integer ierr

call MPI_INIT( ierr )
print *, 'Hello, world!'
call MPI_FINALIZE( ierr )
end
```



All MPI fortran call return an error message

Starting the MPI Environment

- **MPI_INIT ()**

Initializes MPI environment. This function must be called and must be the first MPI function called in a program (exception: **MPI_INITIALIZED**)

Syntax

```
int MPI_Init ( int *argc, char ***argv )
```

```
MPI_INIT ( IERROR )
```

```
INTEGER IERROR
```

NOTE: Both C and Fortran return error codes for all calls.

Exiting the MPI Environment

- **MPI_FINALIZE ()**

Cleans up all MPI state. Once this routine has been called, no MPI routine (even **MPI_INIT**) may be called

Syntax

```
int MPI_Finalize ( );
```

```
MPI_FINALIZE ( IERROR )
```

```
INTEGER IERROR
```

MUST call MPI_FINALIZE when you exit from an MPI program

C and Fortran Language Considerations

- Bindings
 - C
 - All MPI names have an **MPI_** prefix
 - Defined constants are in all capital letters
 - Defined types and functions have one capital letter after the prefix; the remaining letters are lowercase
 - Fortran
 - All MPI names have an **MPI_** prefix
 - No capitalization rules apply
 - last argument is an returned error value

MPI Function Format

- C:

```
#include <mpi.h>
```

```
error = MPI_Xxxxxx(parameter, ...);
```

- Fortran:

```
INCLUDE 'mpif.h'
```

```
CALL MPI_XXXXXX( parameter, ..., IERROR )
```



**don't
forget**

Finding Out About the Environment

- Two important questions that arise early in a parallel program are:
 - How many processes are participating in this computation?
 - Which one am I?
- MPI provides functions to answer these questions:
 - **MPI_Comm_size** reports the number of processes.
 - **MPI_Comm_rank** reports the *rank*, a number between 0 and size-1, identifying the calling process

MPI Rank

- MPI runtime assigns each process a **rank**, which can be used as an ID of the processes
 - ranks start from 0 and extent to N-1
- Processes can perform different tasks and handle different data based on their **rank**

```
...  
if ( rank == 0 ) {  
    ...  
}  
if ( rank == 1 ) {  
    ...  
}  
...
```

Exercise: Hello World

- README.txt
 - Try to answer the questions in the README
 - How is the program compiled?
 - How do you run the parallel program?
- There is a C and Fortran version of the exercise.

Better Hello (C)

```
#include "mpi.h"
#include <stdio.h>

int main( int argc, char *argv[] )
{
    int rank, size;
    MPI_Init( &argc, &argv );
    MPI_Comm_rank( MPI_COMM_WORLD, &rank );
    MPI_Comm_size( MPI_COMM_WORLD, &size );
    printf( "I am %d of %d\n", rank, size );
    MPI_Finalize();
    return 0;
}
```


Better Hello (Fortran)

```
program main
use MPI
integer ierr, rank, size

call MPI_INIT( ierr )
call MPI_COMM_RANK( MPI_COMM_WORLD, rank, ierr )
call MPI_COMM_SIZE( MPI_COMM_WORLD, size, ierr )
print *, 'I am ', rank, ' of ', size
call MPI_FINALIZE( ierr )
end
```

Rank


- The rank identifies different processes within a communicator
- The rank is the basis for any work and data distribution.
- C: `int MPI_Comm_rank(MPI_Comm comm, int *rank)`
- Fortran: `MPI_COMM_RANK(comm, rank, ierror)`
`INTEGER comm, rank, ierror`



```
CALL MPI_COMM_RANK( MPI_COMM_WORLD, myrank, ierror)
```


Some Basic Concepts

- Processes can be collected into *groups*.
- Each message is sent in a *context*, and must be received in the same context.
- A group and context together form a *communicator*.
- A process is identified by its *rank* in the group associated with a communicator.
- There is a default communicator whose group contains all initial processes, called **MPI_COMM_WORLD**.
- Each process has its own number
 - starts with 0
 - ends with (size-1)





Communicator

- Communication in **MPI** takes place with respect to communicators
- **MPI_COMM_WORLD** is one such predefined communicator (something of type “**MPI_COMM**”) and contains group and context information
- **MPI_COMM_RANK** and **MPI_COMM_SIZE** return information based on the communicator passed in as the first argument
- Processes may belong to many different communicators



Split communicator


Split a Large Communicator Into Smaller Communicators



Point to Point communication

MPI Basic Send/Receive


- Basic message passing process. Send data from one process to another



- Questions
 - To whom is data sent?
 - Where is the data?
 - What type of data is sent?
 - How much of data is sent?
 - How does the receiver **identify** it?

MPI Basic Send/Receive

- Data transfer plus synchronization



- Requires co-operation of sender and receiver
- Co-operation not always apparent in code
- Communication and synchronization are combined

Message Organization in MPI

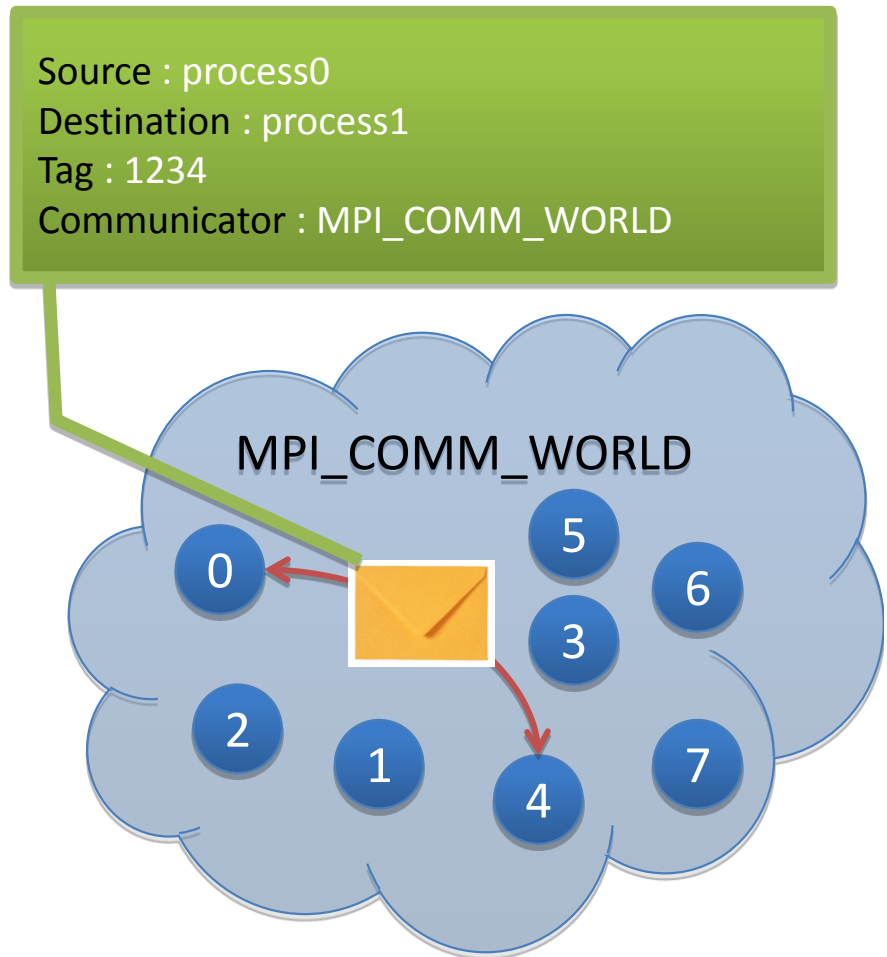
- Message is divided into data and envelope

- data

- buffer
- count
- datatype

- envelope

- process identifier (source/des
- message tag
- communicator



MPI Basic Send/Receive

- Thus the basic (blocking) send has become:

```
MPI_Send ( buf, count, datatype, dest, tag,  
comm )
```

- Blocking means the function does not return until it is safe to reuse the data in buffer. The message may not have been received by the target process.

- And the receive has become:

```
MPI_Recv( buf, count, datatype, source, tag,  
comm, status )
```

- The source, tag, and the count of the message actually received can be retrieved from `status`

MPI C Datatypes

MPI datatype	C datatype
<code>MPI_CHAR</code>	<code>signed char</code>
<code>MPI_SHORT</code>	<code>signed short int</code>
<code>MPI_INT</code>	<code>signed int</code>
<code>MPI_LONG</code>	<code>signed long int</code>
<code>MPI_UNSIGNED_CHAR</code>	<code>unsigned char</code>
<code>MPI_UNSIGNED_SHORT</code>	<code>unsigned short int</code>
<code>MPI_UNSIGNED_LONG</code>	<code>unsigned long_int</code>
<code>MPI_UNSIGNED</code>	<code>unsigned int</code>
<code>MPI_FLOAT</code>	<code>float</code>
<code>MPI_DOUBLE</code>	<code>double</code>
<code>MPI_LONG_DOUBLE</code>	<code>long double</code>
<code>MPI_BYTE</code>	
<code>MPI_PACKED</code>	

MPI Fortran Datatypes

MPI FORTRAN	FORTTRAN datatypes
MPI_INTEGER	INTEGER
MPI_REAL	REAL
MPI_REAL8	REAL*8
MPI_DOUBLE_PRECISION	DOUBLE PRECISION
MPI_COMPLEX	COMPLEX
MPI_LOGICAL	LOGICAL
MPI_CHARACTER	CHARACTER
MPI_BYTE	
MPI_PACKED	

Process Naming and Message Tags

- Naming a process
 - **destination** is specified by (**rank**, **group**)
 - Processes are named according to their rank in the group
 - Groups are defined by their distinct “communicator”
 - **MPI_ANY_SOURCE** wildcard rank permitted in a receive Tags are integer variables or constants used to uniquely identify individual messages
- Tags allow programmers to deal with the arrival of messages in an orderly manner
- MPI tags are guaranteed to range from 0 to 32767 by MPI-1
 - Vendors are free to increase the range in their implementations
- **MPI_ANY_TAG** can be used as a wildcard value

Communication Envelope

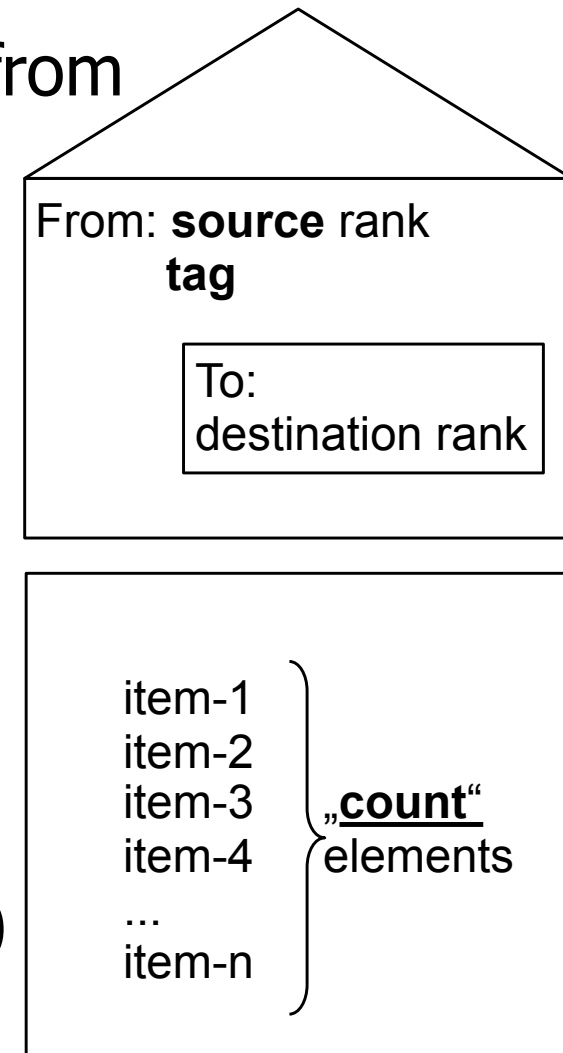
- Envelope information is returned from MPI_RECV in **status**.

- C:

```
status.MPI_SOURCE  
status.MPI_TAG  
count via MPI_Get_count()
```

- Fortran:

```
status(MPI_SOURCE)  
status(MPI_TAG)  
count via MPI_GET_COUNT()
```



Retrieving Further Information

- **Status** is a data structure allocated in the user's program.
- In C:

```
int recvd_tag, recvd_from, recvd_count;
MPI_Status status;
MPI_Recv(..., MPI_ANY_SOURCE, MPI_ANY_TAG, ..., &status )
recvd_tag  = status.MPI_TAG;
recvd_from = status.MPI_SOURCE;
MPI_Get_count( &status, datatype, &recvd_count );
```

- In Fortran:

```
integer recvd_tag, recvd_from, recvd_count
integer status(MPI_STATUS_SIZE)
call MPI_RECV(..., MPI_ANY_SOURCE, MPI_ANY_TAG, .. status, ierr)
tag_recvd  = status(MPI_TAG)
recvd_from = status(MPI_SOURCE)
call MPI_GET_COUNT(status, datatype, recvd_count, ierr)
```

Requirements for Point-to-Point Communications

For a communication to succeed:

- Sender must specify a valid destination rank.
- Receiver must specify a valid source rank.
- The communicator must be the same.
- Tags must match.
- Message datatypes must match.
- Receiver's buffer must be large enough.

Exercise: Send - Recv (1)

Write a simple program where every processor sends data to the next one. You may use as a starting point the basic.c or basic.f90. The program should work as follows:

- Let ntasks be the number of the tasks.
- Every task with a rank less than ntasks-1 sends a message to task myid+1. For example, task 0 sends a message to task 1.
- The message content is an integer array of 100 elements.
- The message tag is the receiver's id number.
- The sender prints out the number of elements it sends and the tag number.
- All tasks with rank ≥ 1 receive messages. You should check the MPI_SOURCE and MPI_TAG fields of the status variable (in Fortran you should check the corresponding array elements). Also check then number of elements that the process received using MPI_Get_count.
- Each receiver prints out the number of elements it received, the message tag, and the rank.
- Write the program using MPI_Send and MPI_Recv

Blocking Communication

- So far we have discussed *blocking* communication
 - **MPI_SEND** does not complete until **buffer** is empty (available for reuse)
 - **MPI_RECV** does not complete until **buffer** is full (available for use)
- A process sending data will be blocked until data in the send buffer is emptied
- A process receiving data will be blocked until the receive buffer is filled
- Completion of communication generally depends on the message size and the system buffer size
- Blocking communication is simple to use but can be prone to deadlocks

Exercise: Send - Recv (2)

- Find out what the program deadlock (.c or .f90) is supposed to do. Run it with two processors and see what happens.
 - a) Why does the program get stuck ?
 - b) Reorder the sends and receives in such a way that there is no deadlock.
 - c) Replace the standard sends with non-blocking sends (MPI_Isend/ MPI_Irecv) to avoid deadlocking. See the man page how to use these non-blocking
 - d) Replace the sends and receives with MPI_SENDRECV.
 - e) In the original program set maxN to 1 and try again.

Sources of Deadlocks

- Send a large message from process 0 to process 1
 - If there is insufficient storage at the destination, the send must wait for the user to provide the memory space (through a receive)
- What happens with

Process 0

Process 1

Send (1)

Send (0)

Recv (1)

Recv (0)

- This is called “unsafe” because it depends on the availability of system buffers.

Some Solutions to the “unsafe” Problem

- Order the operations more carefully:

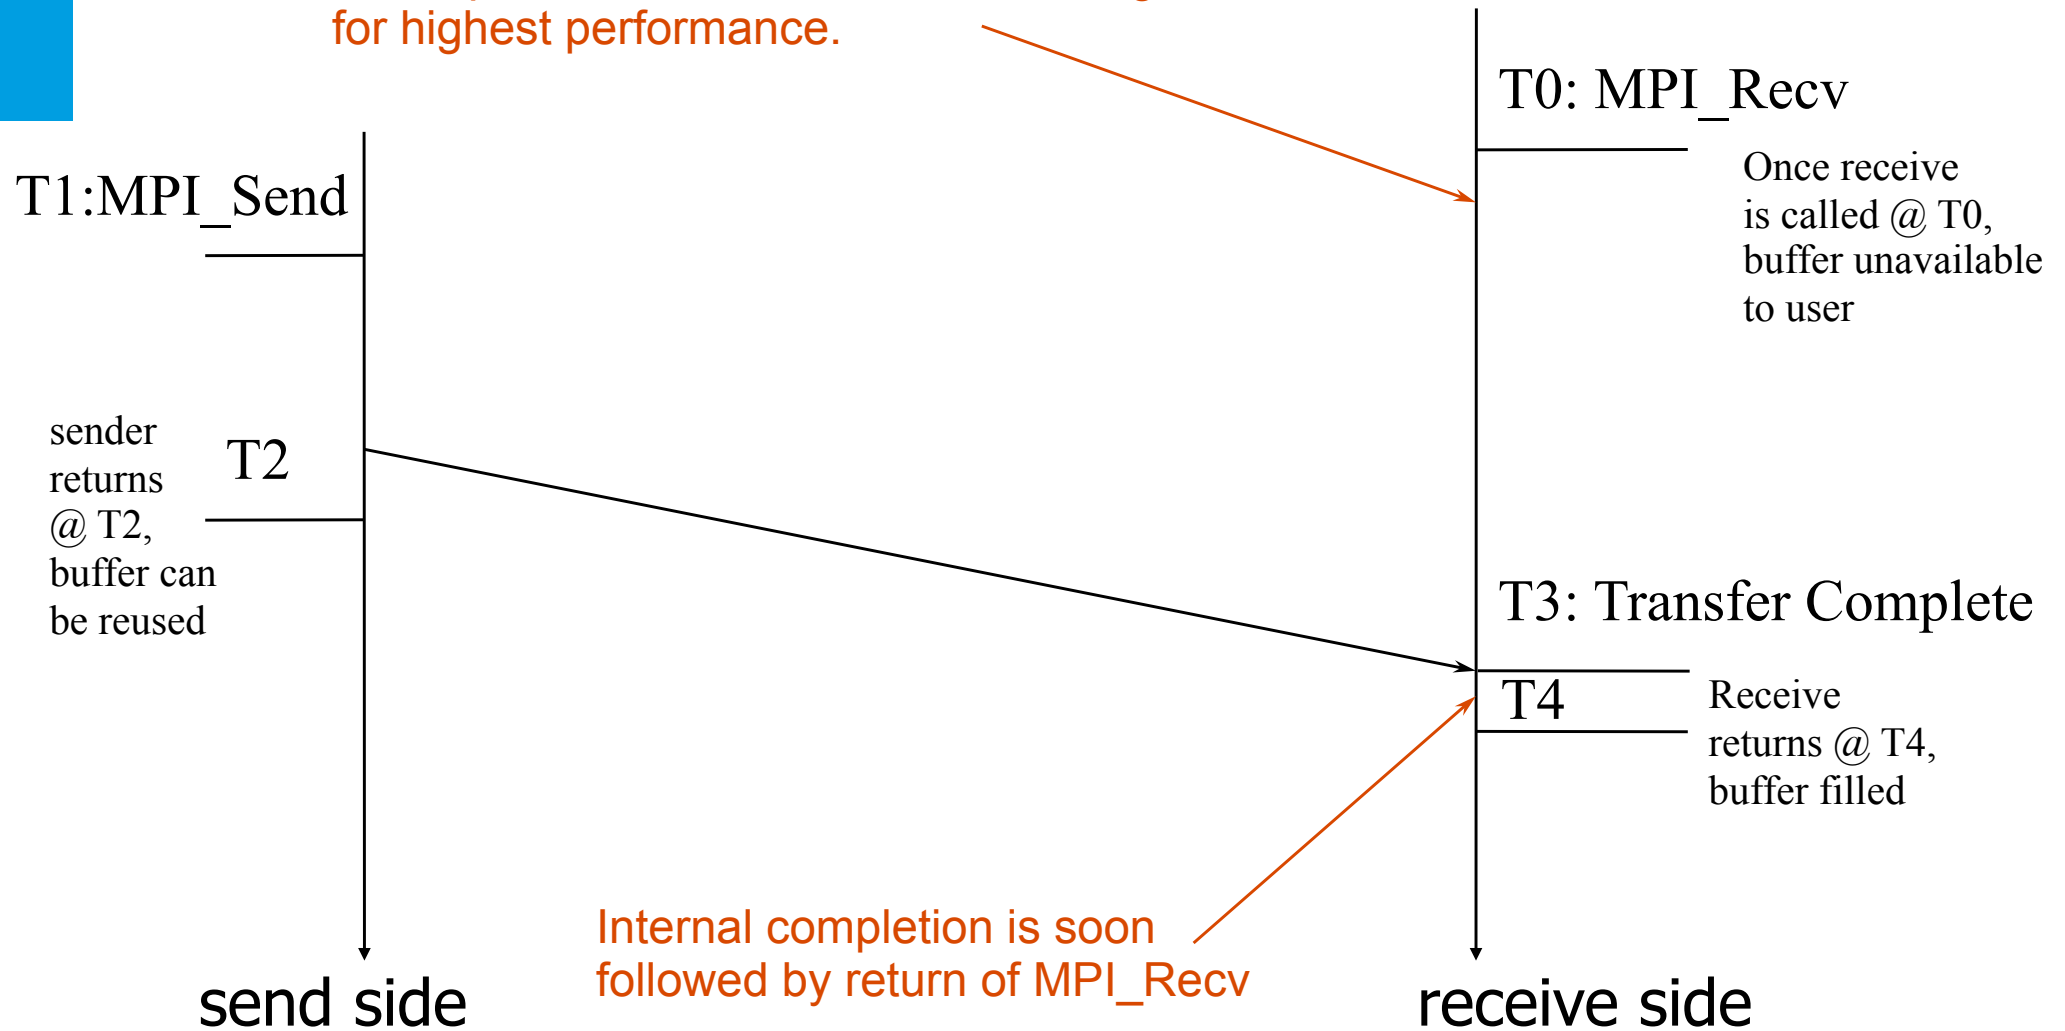
Process 0	Process 1
Send (1)	Recv (0)
Recv (1)	Send (0)

- Use non-blocking operations:

Process 0	Process 1
Isend (1)	Isend (0)
Irecv (1)	Irecv (0)
Waitall	Waitall

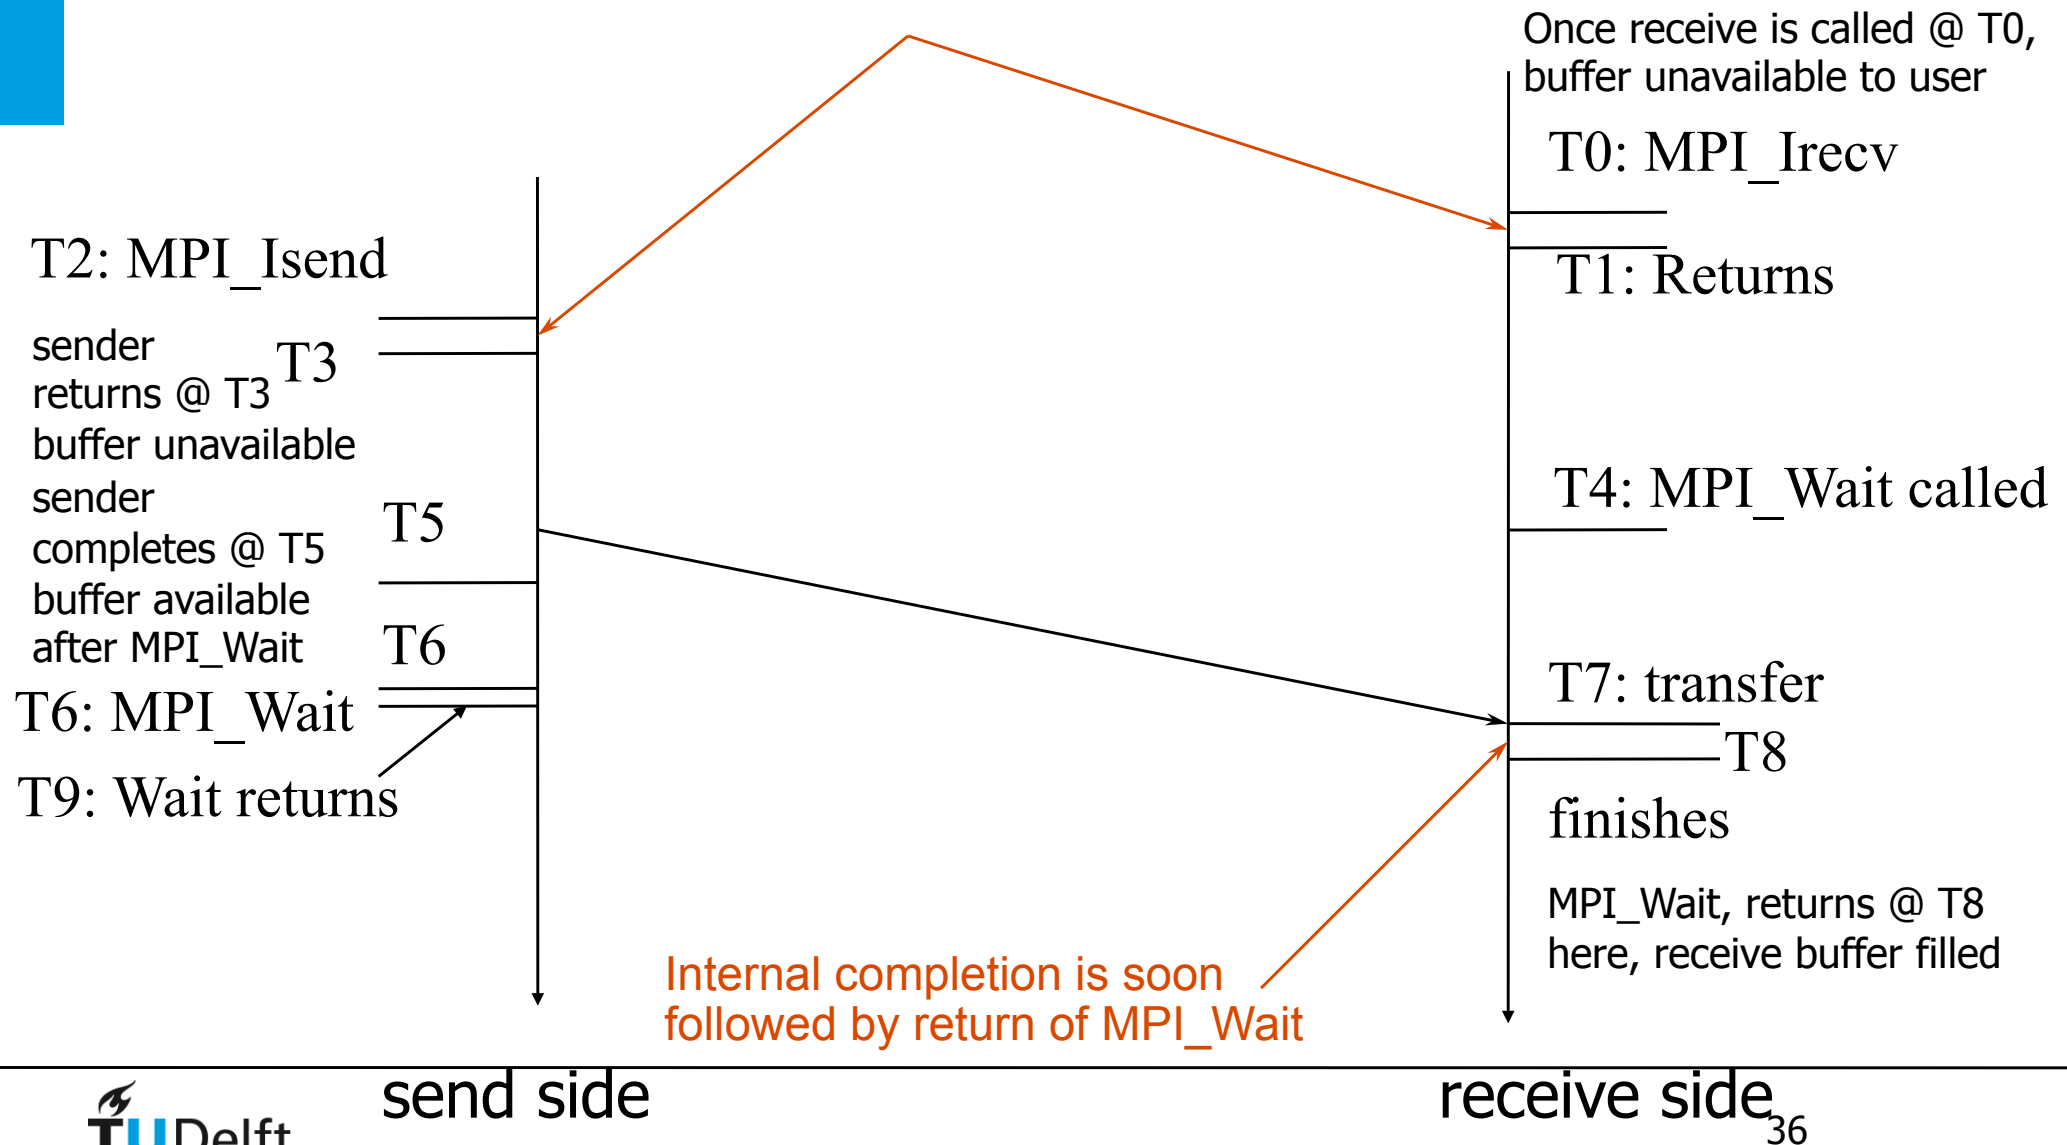
Blocking Send-Receive Diagram (Receive before Send)

It is important to receive before sending,
for highest performance.



Non-Blocking Send-Receive Diagram

High Performance Implementations
Offer Low Overhead for Non-blocking Calls



Non-blocking Receive

- C:

```
MPI_Irecv(buf, count, datatype, source, tag, comm,  
          OUT &request_handle);
```

```
MPI_Wait(INOUT &request_handle, &status);
```

- Fortran:

```
CALL MPI_IRecv (buf, count, datatype, source, tag, comm,  
               OUT request_handle, ierror)
```

```
CALL MPI_WAIT( INOUT request_handle, status, ierror)
```

- buf must **not** be used between Irecv and Wait (in all progr. languages)

Message Completion and Buffering

- A send has completed when the user supplied buffer can be reused

```
*buf = 3;  
MPI_Send ( buf, 1, MPI_INT, ... );  
*buf = 4; /* OK, receiver will always receive 3 */
```

```
*buf = 3;  
MPI_Isend(buf, 1, MPI_INT, ...);  
*buf = 4; /* Undefined whether the receiver will get 3 or 4 */  
MPI_Wait ( ... );
```

- Just because the send completes does not mean that the receive has completed
 - Message may be buffered by the system
 - Message may still be in transit

Non-Blocking Communications

- Separate communication into three phases:
- Initiate non-blocking communication
 - returns Immediately
 - routine name starting with MPI_I...
- Do some work
 - “latency hiding”
- Wait for non-blocking communication to complete

Non-Blocking Communication

- Non-blocking (asynchronous) operations return (immediately) “request handles” that can be waited on and queried

```
MPI_ISEND( start, count, datatype, dest, tag,  
           comm, request )
```


```
MPI_IRecv( start, count, datatype, src, tag,  
           comm, request )
```

```
MPI_WAIT( request, status )
```


- Non-blocking operations allow overlapping computation and communication.
- Anywhere you use `MPI_Send` or `MPI_Recv`, you can use the pair of `MPI_Isend/MPI_Wait` or `MPI_Irecv/MPI_Wait`
- Combinations of blocking and non-blocking sends/receives can be used to synchronize execution instead of barriers

Non-Blocking Examples

- Non-blocking **send**



- Non-blocking **receive**



 = waiting until operation locally completed

Completion

- C:

`MPI_Wait(&request_handle, &status);`

`MPI_Test(&request_handle, &flag, &status);`

- Fortran:

`CALL MPI_WAIT(request_handle, status, ierror)`
completes if request is finished

`CALL MPI_TEST(request_handle, flag, status, ierror)`
test if request_handle is finished without waiting

- one **must** use

- WAIT or
- loop with TEST until request is completed, i.e., `flag == 1` or `.TRUE.`

Multiple Completion's

- It is often desirable to wait on multiple requests
- An example is a worker/manager program, where the manager waits for one or more workers to send it a message

`MPI_WAITALL(count, array_of_requests, array_of_statuses)`

`MPI_WAITANY(count, array_of_requests, index, status)`


`MPI_WAITSOME(incount, array_of_requests, outcount,
array_of_indices, array_of_statuses)`

- There are corresponding versions of **TEST** for each of these

Other Send Modes


- Standard mode (**MPI_Send**, **MPI_Isend**)
 - The standard **MPI Send**, the send will not complete until the send buffer is empty
- Synchronous mode (**MPI_Ssend**, **MPI_Issend**)
 - The send does not complete until after a matching receive has been posted
- Buffered mode (**MPI_Bsend**, **MPI_Ibsend**)
 - User supplied buffer space is used for system buffering
 - The send will complete as soon as the send buffer is copied to the system buffer
- Ready mode (**MPI_Rsend**, **MPI_Irsend**)
 - The send will send eagerly under the assumption that a matching receive has already been posted (an erroneous program otherwise)

Parallel Sum




- Array is originally on process 0 (P0)
- Parallel algorithm
 - Scatter
 - Half of the array is sent to process 1
 - Compute
 - P0 & P1 sum independently their segment
 - Reduction
 - Partial sum on P1 sent to P0
 - P0 sums the partial sums

Parallel Sum




Step 1: Receive operation in scatter

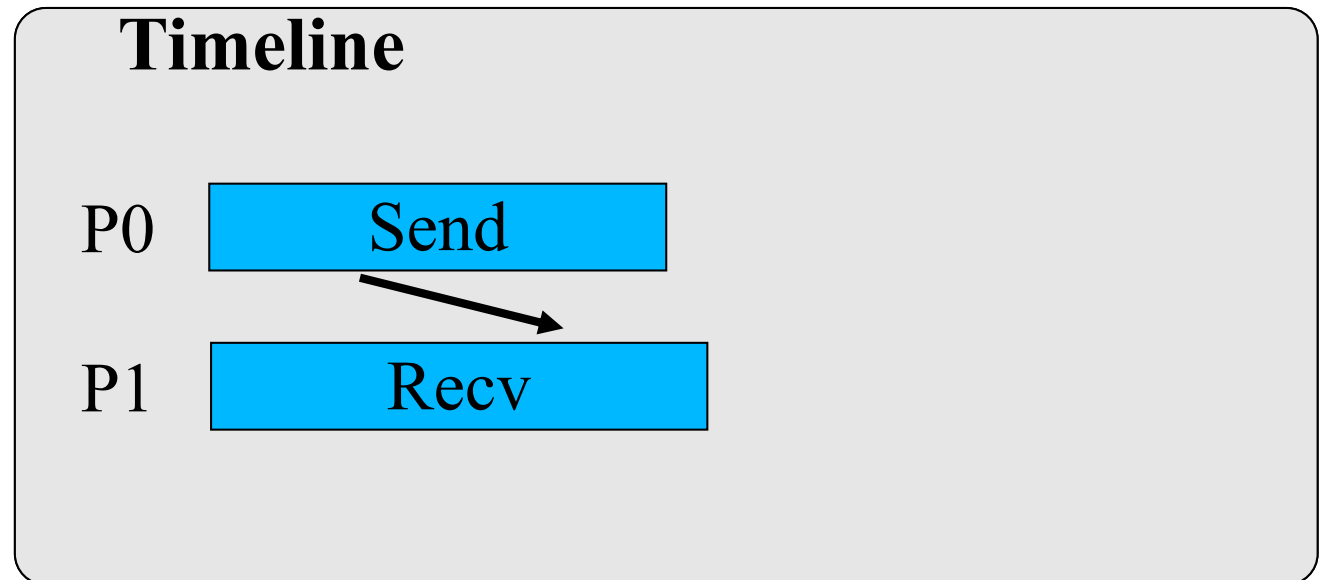


P1 posts a receive to receive half of the array from process 0

Parallel Sum




Step 2: Send operation in scatter

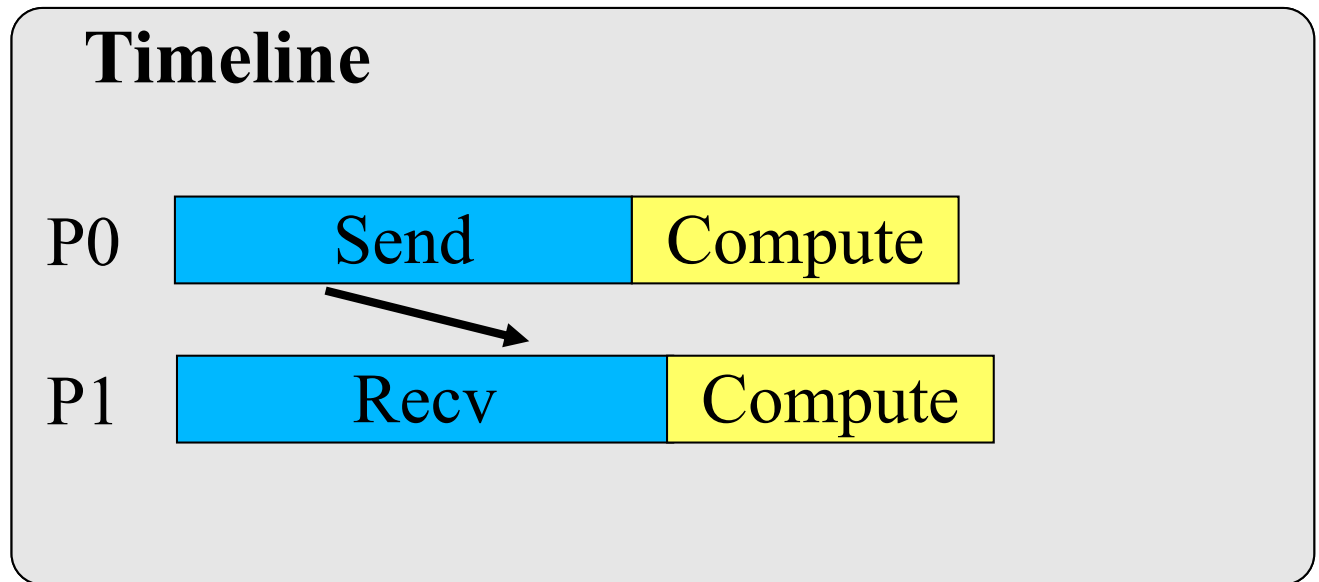


P0 posts a send to send the lower part of the array to P1

Parallel Sum




Step 3: Compute the sum in parallel

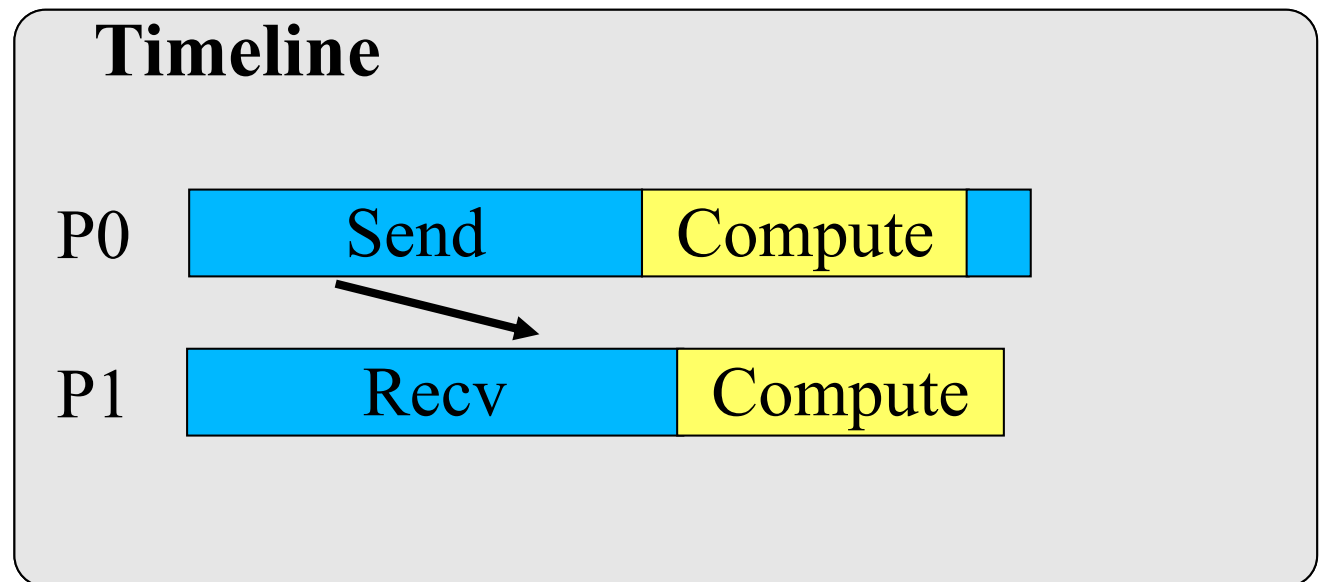


P0 & P1 computes their partial sums and store them locally

Parallel Sum




Step 4: Receive operation in reduction

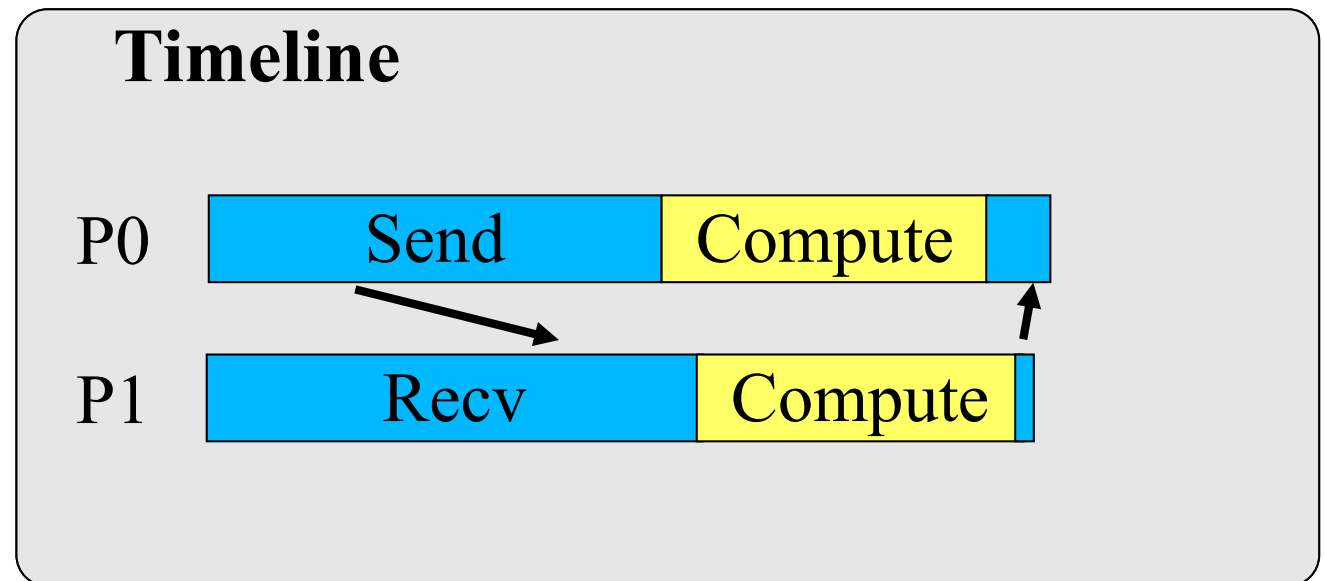


P0 posts a receive to receive partial sum

Parallel Sum




Step 5: Send operation in reduction

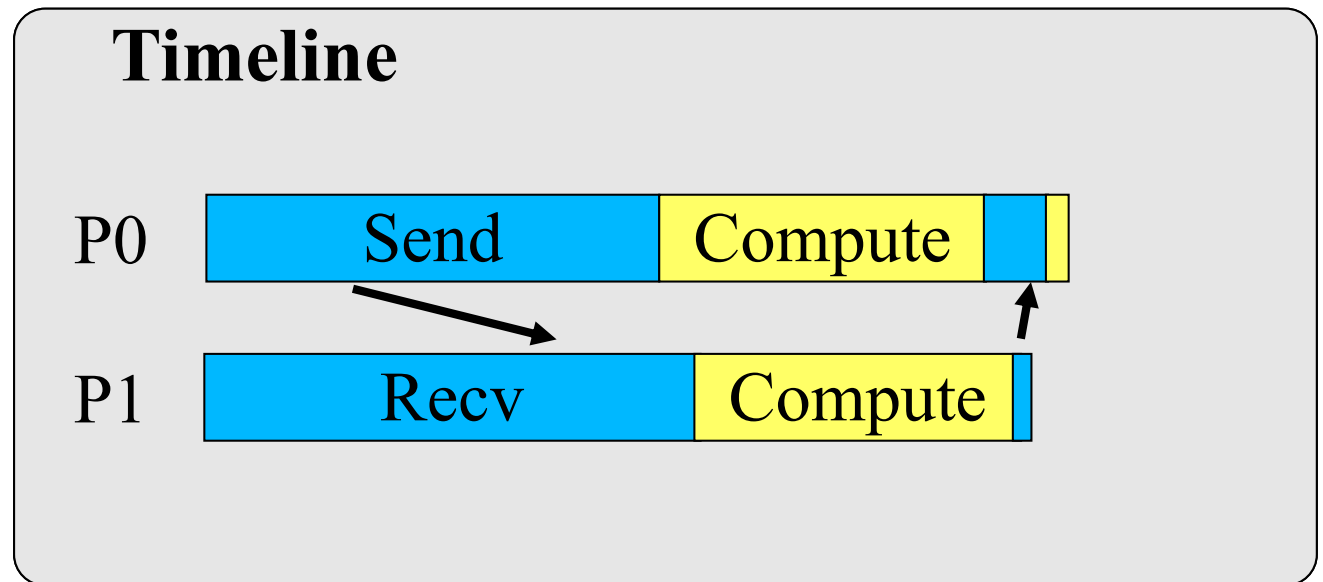


P1 posts a send with partial sum

Parallel Sum



Step 6: Compute final answer



P0 sums the partial sums

Exercise: Parallel sum

```
#include <stdio.h>
#include <stdlib.h>
int main(int argc, char *argv[]){
    int i,N;
    double *array;
    double sum;
    N=100;
    array=malloc(sizeof(double)*N);
    for(i=0;i<N;i++){
        array[i]=1.0;
    }
    sum=0;
    for(i=0;i<N;i++){
        sum+=array[i];
    }
    printf("Sum is %g\n",sum);
}
```

Exercise: Parallel sum

1. Parallelize the sum.c program with MPI
 - The relevant MPI commands can be found back in the README
 - run this program with **two** MPI-tasks
 2. Use MPI_status to get information about the message received
 - print the count of elements received
 3. Using MPI_Probe to find out the message size to be received
 - Allocate an arrays large enough to receive the data
 - call MPI_Recv()
- *_sol.c contains the solution of the exercise.

Probing the Network for Messages

- **MPI_PROBE** and **MPI_IPROBE** allow the user to check for incoming messages without actually receiving them


MPI_PROBE (source, tag, communicator, status)

- **MPI_IPROBE** returns “**flag == TRUE**” if there is a matching message available. **MPI_PROBE** will not return until there is a matching receive available

MPI_IPROBE (source, tag, communicator, flag, status)

Message Order Preservation

- Rule for messages on the same connection, i.e., same communicator, source, and destination rank:
- Messages do not overtake each other.
- This is true even for non-synchronous sends.



- If both receives match both messages, then the order is preserved.

Exercise: Basic Ping-Pong

Ping-pong is a standard test in which two processes repeatedly pass a message back and forth.

Write a program that sends a 'float' array of fixed length, say, ten times back (ping) and forth (pong) to obtain an average time for one ping-pong.

Time the ping-pongs with `MPI_WTIME()` calls.

You may use `pingpong.c` or `pingpong.f90` as a starting point for this exercise.

Investigate how the bandwidth varies with the size of the message.

Basic Ping Pong

rank=0

rank=1


Send (dest=1)
(tag=17)

(tag=23)
Recv (source=1)

Recv (source=0)
Send (dest=0)


```
if (my_rank==0) /* i.e., emulated multiple program */
    MPI_Send( ... dest=1 ...)
    MPI_Recv( ... source=1 ...)
else
    MPI_Recv( ... source=0 ...)
    MPI_Send( ... dest=0 ...)
fi
```

Basic Ping Pong output



Advanced-1 Ping Pong

- A program is written according to the time-line diagram:
 - process 0 sends a message to process 1 (ping)
 - after receiving this message, process 1 sends a message back to process 0 (pong)
- Repeat this ping-pong with a loop of length 50
- Use the timing calls before and after the loop:
- At process 0, print out the transfer time of one message
 - in seconds
 - in μs .
- Use program ping_pong_advanced1.
no need to program yourself.



Advanced-2 Ping-Pong

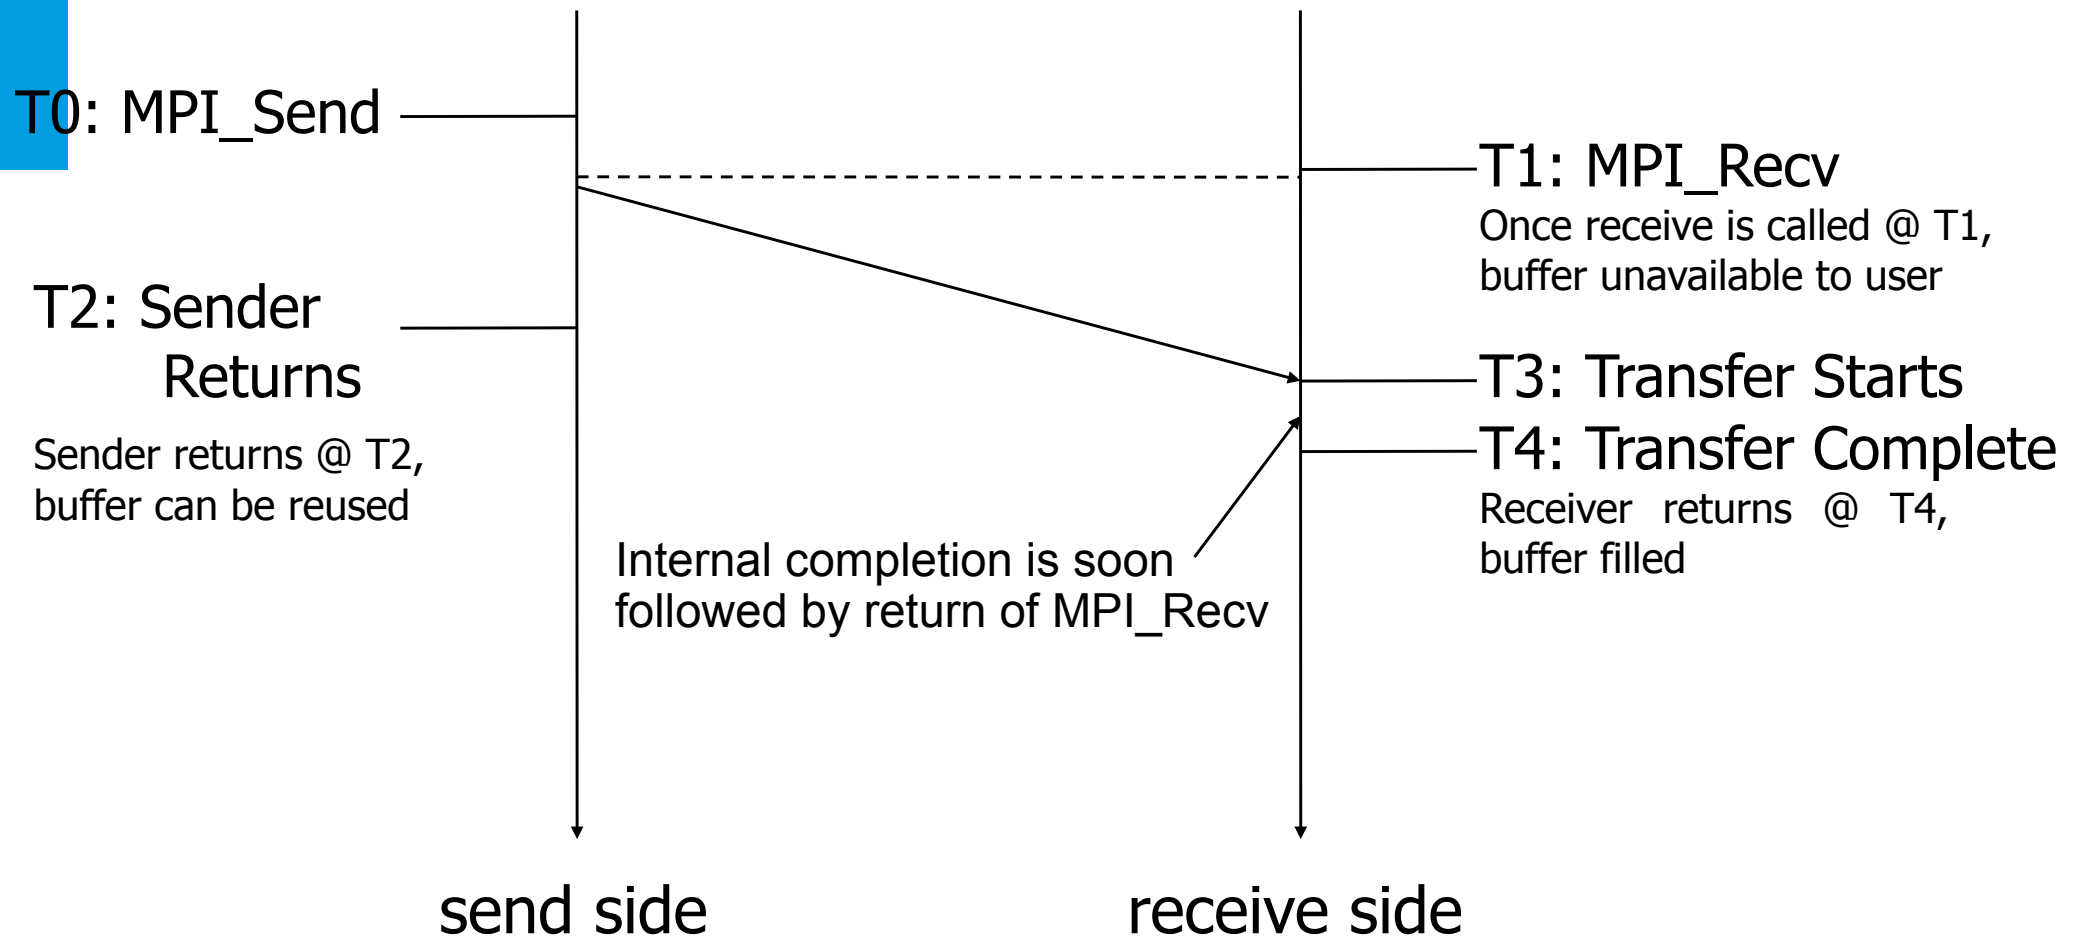
Measure latency and bandwidth

- latency = transfer time for zero length messages
- bandwidth = message size (in bytes) / transfer time
- Print out message transfer time and bandwidth
 - for following send modes:
 - standard send (MPI_Send: ping_pong_advanced2_send.c)
 - synchronous send (MPI_Ssend: ping_pong_advanced2_ssend.c)
 - for following message sizes:
 - 8 bytes (e.g., one double or double precision value)
 - 512 B (= 8*64 bytes)
 - 32 kB (= 8*64**2 bytes)
 - 2 MB (= 8*64**3 bytes)

Standard mode

- Corresponds to the common send functions
 - Blocking: `MPI_Send`
 - Non-blocking: `MPI_Isend`
- It's up to MPI implementation whether communication is buffered or not
- Buffered
 - Can be buffered either locally or remotely
 - The send (blocking) or the completion of a send (non-blocking) may complete before a matching receive
- Non-buffered
 - The send (blocking) or the completion of a send (non-blocking) only complete once it has sent the message to a matching receive
- Standard mode is non-local
 - Successful completion of the send operation may depend on the occurrence of a matching receive.

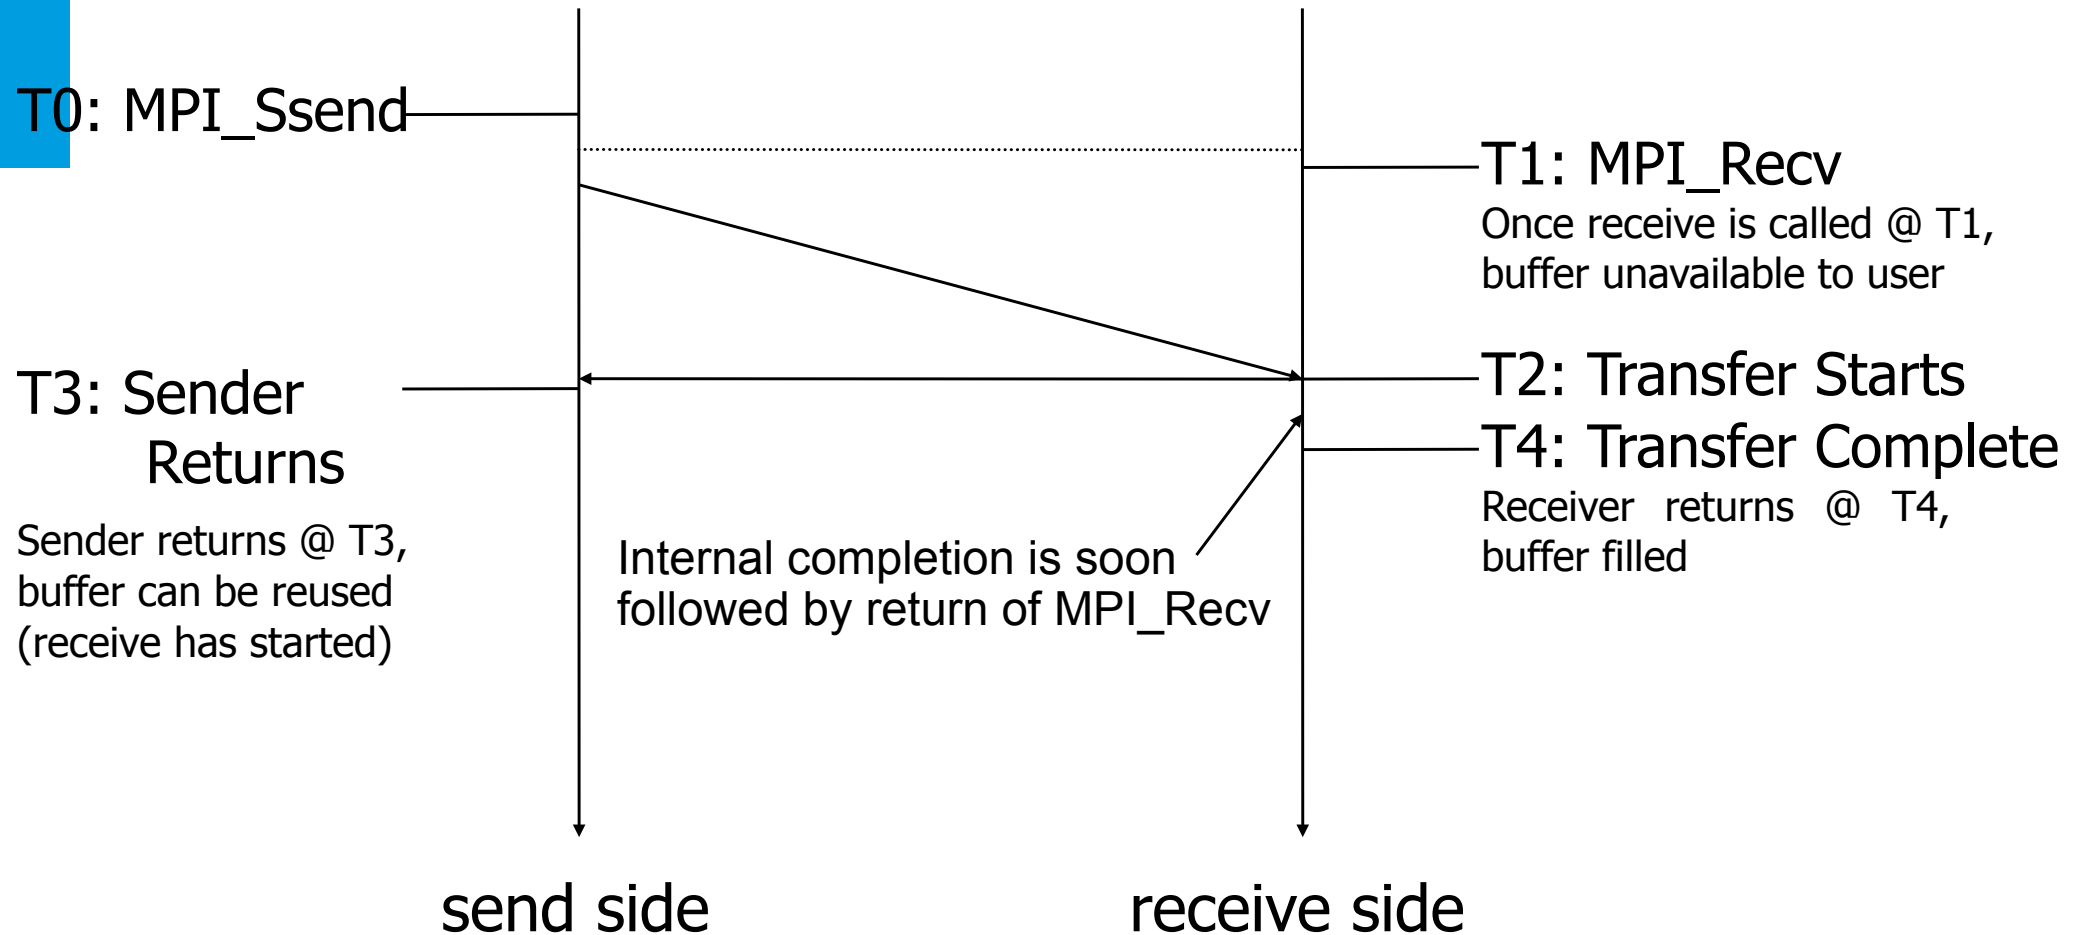
Standard Send-Receive Diagram



Synchronous mode

- Blocking: **MPI_Ssend**
 - Blocking send only returns once the corresponding receive has been posted
 - Same parameters as for standard mode send MPI_Send
- Uses
 - Debugging - potential deadlocks in the program are found by using synchronous sends
 - If many processes send messages to one process its unexpected message buffer can run out if it doesn't pre-post receives. By using MPI_Ssend this can be avoided! Typical example is IO where single process writes data
- Non-blocking: **MPI_Issend**
 - The completion (wait/test) of the send only returns once the corresponding receive has been posted
 - Same parameters as for standard mode send MPI_Isend
 - Useful for debugging - can be used to measure worst case scenario for how long the completion command has to wait


Synchronous Send-Receive Diagram



Exercise Ring


Exercise: Rotating information around a ring

- A set of processes are arranged in a ring.
- Each process stores its rank in MPI_COMM_WORLD into an integer variable `snd_buf`.
- Each process passes this on to its neighbour on the right.
- Each processor calculates the sum of all values.
- Keep passing it around the ring until the value is back where it started, i.e.
- each process calculates sum of all ranks.
- Use non-blocking MPI_Issend
 - to avoid deadlocks
 - to verify the correctness, because blocking synchronous send will cause a deadlock



Non-Blocking Send

- Initiate non-blocking send
→ in the ring example: Initiate non-blocking send to the right neighbor
- Do some work:
→ in the ring example: Receiving the message from left neighbour
- Now, the message transfer can be completed
- Wait for non-blocking send to complete /



Non-Blocking Receive

- Initiate non-blocking receive
 —→ in the ring example: Initiate non-blocking receive from left neighbor
- Do some work:
 —→ in the ring example: Sending the message to the right neighbor
- Now, the message transfer can be completed
- Wait for non-blocking receive to complete

