



WWW.TACC.UTEXAS.EDU



OpenMP Tasking Concepts

TACC: Advanced OpenMP

PCSE: class: 4/17/18

PRESENTED BY:

Kent Milfeld, Lars Koesterke, Charlie Dey

Tasking

Review Worksharing and Limitations

Basic Task Syntax and Operations

Task Synchronization

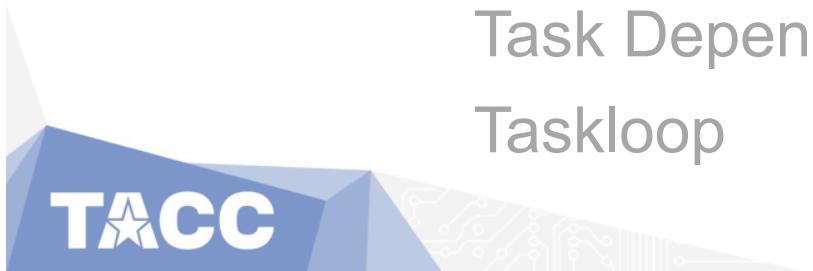
Running Tasks in Parallel

Data-sharing and firstprivate Default for Tasks

Common Use Cases for Tasks

Task Dependencies

Taskloop



Worksharing

- TEAM OF THREADS does work
- LOOP ITERATION PARTITIONING (chunks of independent work)
- DATA ENVIRONMENT is set by parallel region, altered by clauses.
- IMPLIED BARRIER forces threads to wait at end.

```
omp parallel  
{  
}  
→  
omp parallel  
{  
    <Worksharing Loops>  
}
```

Example - schedule(static,16), threads = 4

```
!$omp parallel do schedule(static,16)
    do i=1,128
        A(i)=B(i)+C(i)
    enddo
```

<u>thread0</u> :	do i=1,16 A(i)=B(i)+C(i) enddo do i=65,80 A(i)=B(i)+C(i) enddo
<u>thread1</u> :	do i=17,32 A(i)=B(i)+C(i) enddo do i = 81,96 A(i)=B(i)+C(i) enddo
<u>thread2</u> :	do i=33,48 A(i)=B(i)+C(i) enddo do i = 97,112 A(i)=B(i)+C(i) enddo
<u>thread3</u> :	do i=49,64 A(i)=B(i)+C(i) enddo do i = 113,128 A(i)=B(i)+C(i) enddo

Worksharing Limitation

- Worksharing requires ability to know number of work units before execution (loop count).
- Dynamic scheduling of parallel loop work-chunks is limited. Think of work chunks being queued to execute from a single FIFO queue*.
- Ideal for “data parallel”, but “task parallel” requires inconvenient construction.

*This can be somewhat limited, but reasonable for monotonically increasing/decreasing work per iteration.



Learning Objective

- Tasking --
 - Review Worksharing and Limitations
 - Basic Task Syntax and Operations
 - Task Synchronization
 - Running Tasks in Parallel
 - Data-sharing and firstprivate Default for Tasks
 - Common Use Cases for Tasks
 - Task Dependencies
 - Taskloop



Getting Started with Tasking

Easy to begin using, for a novice

Uncomplicated syntax and coding for many cases

Task clauses provide enough flexibility to be efficient (in programming syntax and execution) for more complicated use cases.

Is it easy? – that depends.
It is somewhat like worksharing—
There are simple and complex cases.



What is Tasking for?

Irregular Computing:

While loop, execute independent iterations in parallel

Follow pointers in list until a NULL pointer is reached,
performing independent work for each pointer position.

Note: the pointer chase is inherently serial but if work at each pointer position is independent, then work can be executed in parallel.

Follow nodes in tree graph & perform independent work at nodes

Ordered executions that have task (work) with dependences

Generating a task

We first go over generating tasks in a serial region – computationally impractical.

But...it is a practical base for developing concepts.

Hold on– we'll get to executing tasks concurrently in a parallel region!



Generating a task

task is a directive, with clauses (not shown)

- A task is “deferrable”. It can be executed **immediately** or **later**.
- Deferred tasks are queued.

Tasks in
a serial
region.

Tasks do
independent
work.

C/C++

```
#pragma omp task
    foo(j);

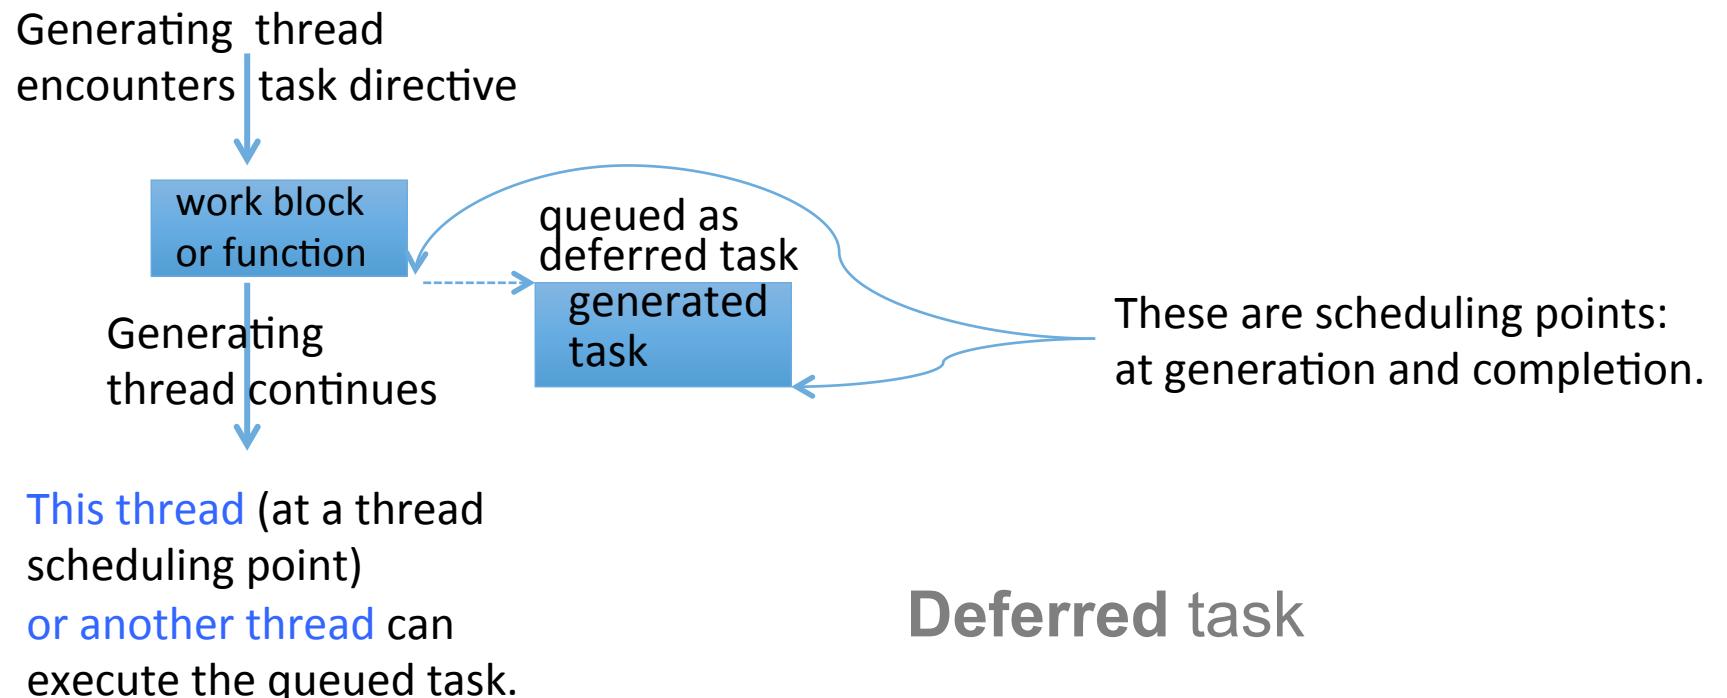
#pragma omp task
{
    for(i=0;i<n;i++) {...};
}
```

F90

```
!$omp task
    foo(j)
 !$omp end task

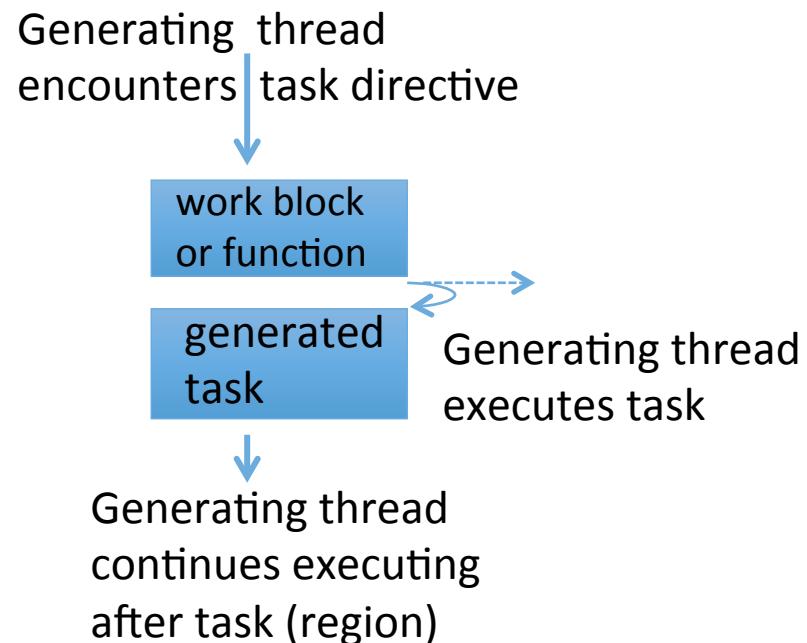
 !$omp task
    do i = 1,n; ... ;enddo
 !$omp end task
```

Deferred Task



Deferred task

Immediate Task

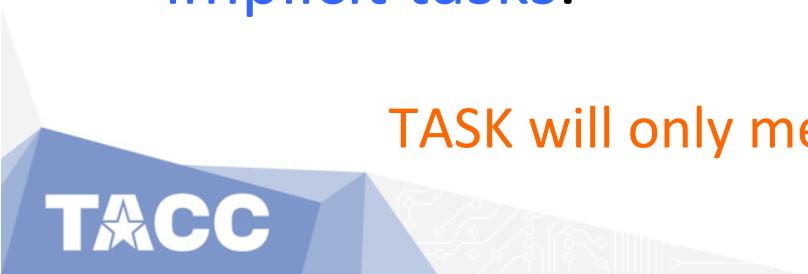


Immediate task

What is a Task

- It is work generated by a **task construct** – yeah, yeah we know that now!
IT IS an explicit task.
- Threads of a **parallel** region, executing the replicated work-- these are also tasks, with the distinctive name:
implicit tasks.

TASK will only mean Explicit Task unless otherwise stated.



Learning Objective

- Tasking --
 - Review Worksharing and Limitations
 - Basic Task Syntax and Operations
 - Task Synchronization**
 - Running Tasks in Parallel
 - Data-sharing and firstprivate Default for Tasks
 - Common Use Cases for Tasks
 - Task Dependencies
 - Taskloop



Synchronizing tasks (sibling tasks)

If a task is deferred, use the **taskwait** construct to wait for completion at some point in the code.

C/C++

```
#pragma omp task
{ foo(j); }

#pragma omp task
{ for(i=0;i<n;i++) {...}; }

#pragma omp taskwait
```

F90

```
!$omp task
call foo(j)
!$omp end task

!$omp task
do i = 1,n; ... ;enddo
!$omp end task
!$omp taskwait
```

Synchronizing tasks (nested tasks)

A **taskgroup** construct waits for all sibling and their descendants.

Nested Tasks

```
#pragma taskgroup C/C++  
{  
    #pragma omp task  
    foo(j);  
  
    #pragma omp task  
    { for(i=0;i<n;i++)  
        #pragma omp task  
        foo(i);  
    }  
}
```

```
!omp taskgroup F90  
!  
omp task  
call foo(j)  
!  
omp end task  
  
!  
omp task  
do i = 1,n  
!  
omp task  
call foo(i);  
!  
omp end task  
enddo  
!  
omp end task  
!  
omp end taskgroup
```

Tasking

Review Worksharing and Limitations

Basic Task Syntax and Operations

Task Synchronization

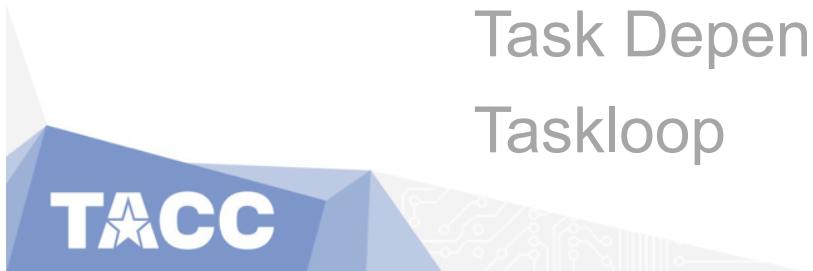
Running Tasks in Parallel

Data-sharing and firstprivate Default for Tasks

Common Use Cases for Tasks

Task Dependencies

Taskloop



Generating Concurrent Tasks—in a parallel region

First, create a **team of threads**, to work on tasks.

Use a single thread to **generate tasks**.

C/C++

```
#pragma omp parallel num_threads(4)
{
    #pragma omp single
    {
        //generate multiple tasks
    }
}
```

F90

```
!$omp parallel num_threads(4)

 !$omp single

 !generate multiple tasks

 !$omp end single
 !$omp end parallel
```



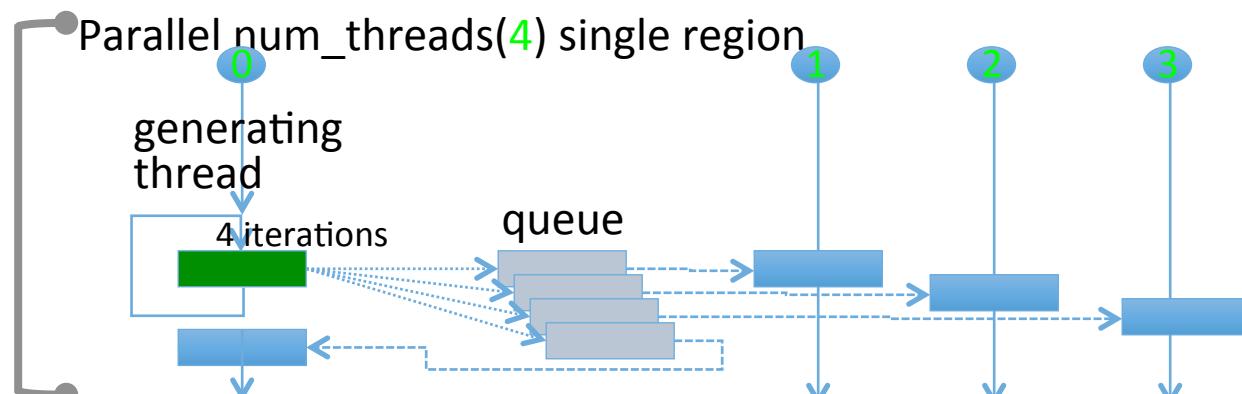
Generate multiple tasks: with loop (while/do/for) or recursion

Tasks in parallel region

Threads of Team will dequeue & execute tasks

Shared variables of parallel region are also shared by tasks

Tasks obey explicit and implicit barriers, also taskwait & taskgroup



Tasking

Review Worksharing and Limitations

Basic Task Syntax and Operations

Task Synchronization

Running Tasks in Parallel

Data-sharing and firstprivate Default for Tasks

Common Use Cases for Tasks

Task Dependences

Taskloop



Tasks: Basic Data-sharing Attributes

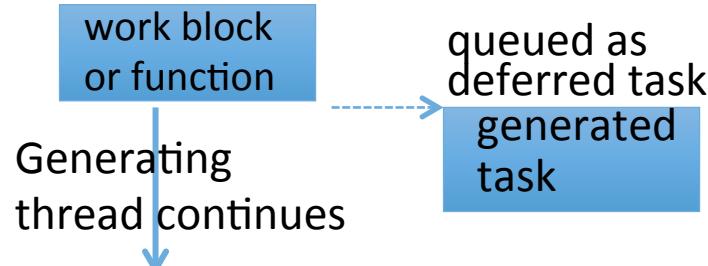
- If the task generating construct is in a parallel region any shared variables remain shared.
- for/do index variables of a worksharing loop are private (see spec.)
- private variables in an enclosing construct are firstprivate for the task.

The Point: The index passed to a task needs to be first private.



Deferred Task

Generating thread
encounters task directive



```
#pragma omp task firstprivate(j)  
foo(j);  
  
j++;
```

Basic Concept:
The argument value at generation time
is needed— not later when it is run.

```
!$omp task firstprivate(j)  
foo(j)  
!$omp end task  
  
j = j+1
```

In a parallel region j is shared, and hence needs to be declared firstprivate.

Scheduling Optimization

For a small number of threads and tasks, and a large diversity in task work—an imbalance will occur. Even with moderate diversity and large thread and task counts, an imbalance may be present. The **priority clause** can alleviate this problem. (Larger # = higher priority.)

```
for (i=0;i<N; i++) {  
    #pragma omp task priority(i+1)  
    compute_array(array[i], N);  
}
```

```
do i=1,N  
    !$omp task priority(i)  
    call compute_array(matrix(:, i), N)  
    !$omp end task  
enddo
```

Tasking

Review Worksharing and Limitations

Basic Task Syntax and Operations

Task Synchronization

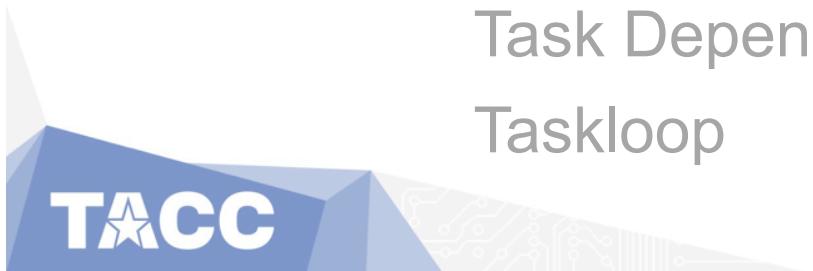
Running Tasks in Parallel

Data-sharing and `firstprivate` Default for Tasks

Common Use Cases for Tasks

Task Dependences

Taskloop



While loop

- `firstprivate` clause is necessary, since `cntr` is shared and value must be captured for work.

```
int cntr = 100;
#pragma omp parallel
#pragma omp single

while(cntr>0){

#pragma omp task firstprivate(cntr)
printf("cntr=%d\n",cntr);

    cntr--;
}
```

```
integer cntr = 100
!$omp parallel
!$omp single

do while(cntr>0)

    !$omp task firstprivate(cntr)
        print*, "cntr= ",cntr
    !$omp end task

    cntr = cntr - 1
enddo
 !$omp end single
 !$omp end parallel
```

Exploiting tasks within while loop

The generating loop is executed SERIALLY, but concurrently with the dequeued tasks.

- So, the non-tasking loop parts should not be costly.
- Any generated tasks can be picked up directly by other team members.

```
#pragma omp single
while(cntr>0) {
    #pragma omp task firstprivate(cntr)
        printf("cntr=%d\n",cntr);

    cntr--;
}
```

The diagram illustrates the execution flow of the code. It shows three parallel regions defined by brackets on the right side:

- Serial – generation:** The outermost region, containing the entire while loop.
- Parallel – Tasking:** The region where tasks are generated and executed.
- Serial – incrementer:** The innermost region, where the counter is decremented.

Pointer Chasing

- *ptr* points to a C/C++ structure or F90 defined type

```
int *ptr;
...//initialize pointer
#pragma omp parallel
#pragma omp single

while(ptr) {
    #pragma omp task firstprivate(ptr)
    process(ptr);

    ptr = ptr->next;
}
```

```
integer,pointer :: ptr
...! initialize pointer
!$omp parallel
!$omp single

do while(associated(ptr))
    !$omp task firstprivate(ptr)
        process(ptr)
    !$omp end task

    ptr = ptr%next
enddo

!$omp end single
!$omp end parallel
```

Using Recursion in C/C++

- Starting point for tree searches.
- Note, the recursion function, not the “process”, is directly tasked.

```
#pragma omp parallel
#pragma omp single
    chase(ptr);

...
void chase(node *ptr) {

    if(ptr->next) {
        #pragma omp task
        chase(ptr->next);
    }
    process(ptr)
}
```

Just another way
to build a queue
of tasks— but with
recursion (bunch
of context switching)

Undeferred Tasks with **if** clause

```
while(ptr) {  
  
    usec=ptr->cost*factor;  
  
    #pragma omp task if(usec>0.01) firstprivate(ptr)  
    process(ptr)  
  
    ptr = ptr->next;  
}
```

If the **if** argument is false (exec time (usec) is less than 0.01), task is undeferred.

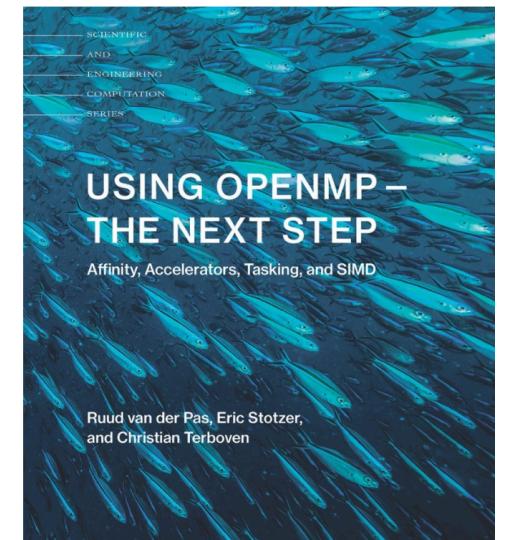
- Generating thread will suspend generation
- Generating thread will execute the task
- Generating thread will resume generation

--The END--

Questions?

References:

OpenMP Programming: The Next Step



More Steps? We can cover Task Dependences, time permitting.

Tasking

Review Worksharing and Limitations

Basic Task Syntax and Operations

Task Synchronization

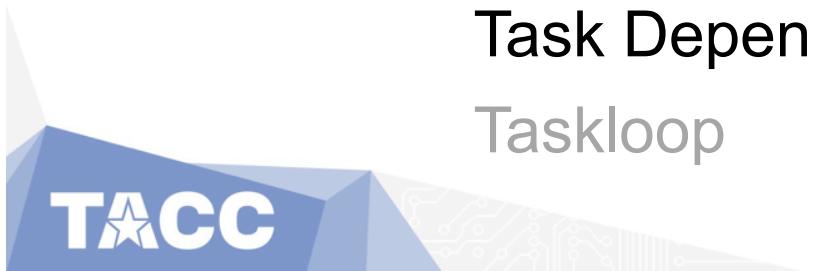
Running Tasks in Parallel

Data-sharing and firstprivate Default for Tasks

Common Use Cases for Tasks

Task Dependences

Taskloop



Depend clause

Syntax: ... **task depend(** *dependence-type*: *list* **)**

- Dependences are derived from *dependence-type* and the *list* items of the depend clause.

dependence-type: == what the task needs to do with *list* item

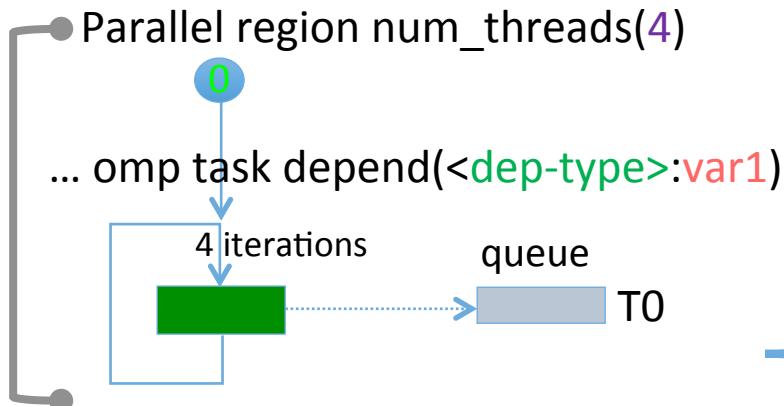
in think of as a **read**

out think of as a **write**

inout think of as a **read and then a write**

list item (variable) needing to “Read” or “Write”

Dependences



Task 1 execution dependence

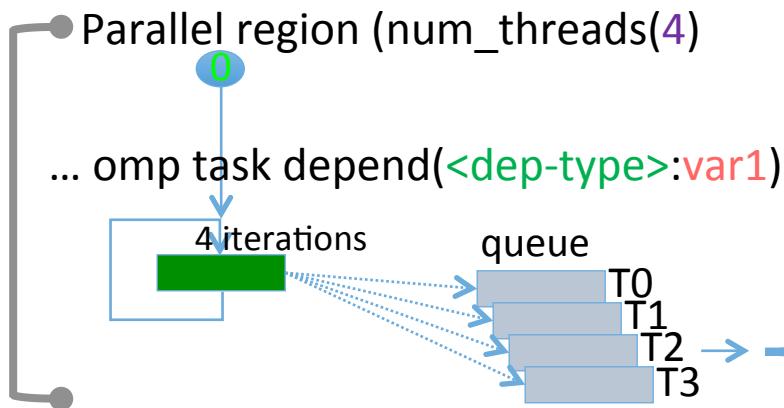
When runtime executes `T0`, it checks previously generated tasks for identical list items (`var1`).

There are no previous tasks— so this task has NO dependence. EVEN if it has an IN (read) dependence type!



Dependences

Task 3 execution dependence



T3 checks previously generated tasks for identical list times (var1).

If an identical list item exists in previously generated tasks, T3 adheres to the *dependence-type*.



Task depend clause

Flow Control (RaW, Read after Write)

```
x = 1;  
#pragma omp parallel  
#pragma omp single  
{  
    #pragma omp task shared(x) depend(out: x)  
    x = 2;  
    #pragma omp task shared(x) depend( in: x)  
    printf("x = %d\n", x);  
}  
...  
...
```

T1 is put on queue,
sees no previously
queued tasks
with x identifier →
No dependences.

T2 is put on queue,
sees previously queued
task with identifier →
Has RaW dependence.

Print value is
always 2.

Task depend clause

Anti-dependence (WaR, write after read)

```
x = 1;  
#pragma omp parallel  
#pragma omp single  
{  
    #pragma omp task shared(x) depend( in: x)  
    printf("x = %d\n", x);  
    #pragma omp task shared(x) depend(out: x)  
    x = 2;  
}  
....
```

T1 is put on queue,
sees no previously
queued tasks
with x identifier →
No dependences.

T2 is put on queue,
sees previously queued
task with identifier →
has WaR dependence.

Print value is
always 1.

Task depend clause

Output Dependence (WaW, Write after Write)

```
x = 1;  
#pragma omp parallel  
#pragma omp single  
{  
    #pragma omp task shared(x) depend(out: x)  
    printf("x = %d\n", x);  
    #pragma omp task shared(x) depend(out: x)  
    x = 2;  
}
```

T1 is put on queue,
sees no previously
queued tasks
with x identifier →
No dependences.

T2 is put on queue,
sees previously queued
task with identifier →
has WaW dependence.

Print value is
always 1.

Task depend clause

(RaR, no dependence)

```
x = 1;  
#pragma omp parallel  
#pragma omp single  
{  
    #pragma omp task shared(x) depend(in: x)  
    printf("x = %d\n", x);  
    #pragma omp task shared(x) depend(in: x)  
    x = 2;  
}
```

T1 is put on queue,
sees no previously
queued tasks
with x identifier →
No dependences.

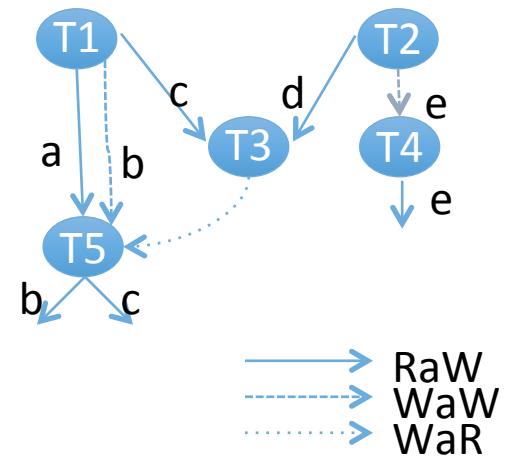
T2 is put on queue,
sees previously queued
task with x identifier →
has NO ordering
(because it is RAR)

Print value is
1 or 2.

Task depend clause

Following a graph

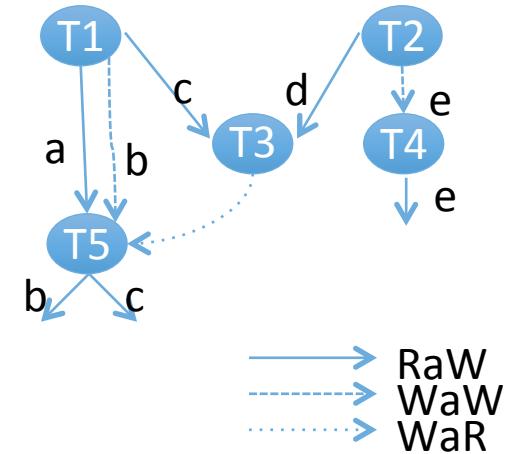
```
#pragma omp parallel
#pragma omp single
{
T1 #pragma omp task depend(out:a,b,c)
    f1(&a, &b, &c);
T2 #pragma omp task depend(out:d,e)
    f2(&d, &e);
T3 #pragma omp task depend(in:c,d)
    f3(c,d);
T4 #pragma omp task depend(out,e)
    f4(&e);
T5 #pragma omp task depend(in:a) depend(out:b,c)
    f5(a,&b,&c)
}
```



Task Depend Clause

Following non-computed variables-- works, too.

```
#pragma omp parallel
#pragma omp single
{
T1 #pragma omp task depend(out:t1,t2,t3)
    f1(&a, &b, &c);
T2 #pragma omp task depend(out:t4,t5)
    f2(&d, &e);
T3 #pragma omp task depend(in:t3,t4)
    f3(c,d);
T4 #pragma omp task depend(out,t5)
    f4(&e);
T5 #pragma omp task depend(in:t1)depend(out:t2,t3)
    f5(a,&b,&c)
}
```



Tasking

Review Worksharing and Limitations

Basic Task Syntax and Operations

Task Synchronization

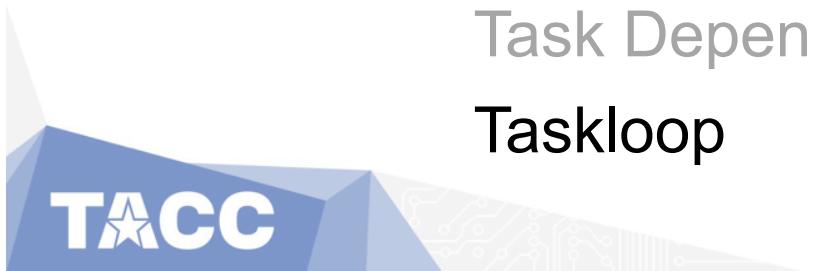
Running Tasks in Parallel

Data-sharing and firstprivate Default for Tasks

Common Use Cases for Tasks

Task Dependencies

Taskloop



taskloop

Iterations of loops are executed as tasks (of a taskgroup)

Single generator needed

All team members are not necessary

Implied taskgroup for the construct

Syntax: . . . omp taskloop [*clauses*]

- some clauses:

grainsize
numtasks or

number of iterations assigned to a task (See spec 4 details.)
number of tasks to be executed concurrently (See spec 4 details.)

default: --number of tasks & iterations/task implementation defined

untied

tasks need not be continued by initial thread of task

nogroups

don't create a task group

priority

for each task (default 0)

Taskloop

- ptr points to a structure in C/C++, defined type in F90

```
void parallel_work(void) { // execute by single in parallel
    int i, j;
    #pragma omp taskgroup
    {
        #pragma omp task
        long_running(); // can execute concurrently

        #pragma omp taskloop private(j) grainsize(500) nogroup
        for (i = 0; i < 10000; i++) //can execute concurrently
            for (j = 0; j < i; j++)
                loop_body(i, j);

    } // end taskgroup
}
```

Summary

- Tasks are used mainly in irregular computing.
- Tasks are often generated by a single thread.
- Task generation can be recursive.
- Depend clause can prescribe dependence.
- Priority provides hint for execution order.
- First private is default data-sharing attribute, shared variables remain shared.
- Untied generator task can assure generation progress.