



Parallel Computing for Science & Engineering



Intel Xeon Phi,

02/07/2018

Instructors:
Lars Koesterke, TACC
Charlie Dey, TACC



Intel Teraflop chip, 2007

OpenMP Setting Number of Threads

1. *Before run: by Environment Variable, OMP_NUM_THREADS*
2. *In an OpenMP API Function*
3. *In a Clause:*

1.	<pre>% setenv OMP_NUM_THREADS # % export OMP_NUM_THREADS=#</pre>	Env. Var.	precedence over implementation defined.
2.	<pre>omp_set_num_threads(#)</pre>	API	precedence over 1.
3.	<pre>!\$omp parallel num_threads(int. expression) #pragma omp parallel num_threads(int expression)</pre>	Clause	precedence over 1 and 2.

OpenMP Setting Number of Threads

**0. `OMP_THREAD_LIMIT` (Env.Var.) Sets limit over all methods.
Default value = implementation defined (e.g 2^{32} on intel)**

1. `% setenv OMP_NUM_THREADS #`
`% export OMP_NUM_THREADS=#`

2. `omp_set_num_threads(#)`
`omp_set_num_threads(omp_get_num_procs\(\)) *`

3. `!$omp parallel num_threads(int. expression)`
`#pragma omp parallel num_threads(int expression)`

[omp_get_num_threads\(\)](#)

Reports # of threads in team.

`omp_get_max_threads()`

Reports # of threads by 1, 2 **.

`export OMP_THREAD_LIMIT=#`

If 1,2 or 3 exceeds limit#, use limit #

`omp_get_thread_limit()`

Reports `OMP_THREAD_LIMIT`.



* If hyperthreading (2-SMT) is on : `procs=cores*2`

** "Max Threads is # of threads that could be used by a parallel region without a `num_threads` clause."

OpenMP Setting Number of Threads

```
#include <stdio.h>
#include <omp.h>
//env. Var. OMP_NUM_THREADS =12
int main(){
    int nt, N=1000;
    printf("PROCS = %d\n",\
        omp_get_num_procs());

    omp_set_num_threads(4);

    #pragma omp parallel num_threads(N/100)
    {

        printf("thrds %d max %d\n",\
            omp_get_num_threads(),\
            omp_get_max_threads());
    }
}
```

```
program nthreads

! Env. var. OMP_NUM_THREADS=12
use omp_lib
integer, parameter :: N=1000
print*, "PROCS = ", &
        omp_get_num_procs()

call omp_set_num_threads(4)

!$omp parallel num_threads(N/100)

    print*, &
        "thrds", omp_get_num_threads(), &
        "max",  omp_get_max_threads()
!$omp end parallel
end program
```



lscpu command reports →

Thread(s) per core: 2
Core(s) per socket: 12
Socket(s): 2

OpenMP Thread and Memory Location

Where do **threads**/processes and **memory allocations** go?

Default: Decided by policy when **process exec' d or thread forked**, and **on write to memory** . Processes and threads can be rescheduled to different sockets and cores.

Ways Process Affinity and Memory Policy can be changed:

- 1.) Dynamically on a running process (knowing process id)
- 2.) At process execution (with wrapper command)
- 3.) Within program through F90/C API

Barrier Example

- Barrier: Each thread waits until all threads arrive

```
#pragma omp parallel shared (A, B, C) private(id)
{
    id=omp_get_thread_num();
    A[id] = big_calc1(id);
    #pragma omp barrier
    #pragma omp for
    for(i=0;i<N;i++){
        C[i]=big_calc3(i,A,N);
    }←————— Implicit barrier
    #pragma omp for nowait
    for(i=0;i<N;i++){
        B[i]=big_calc2(C, I, N);
    }←←————— No Implicit barrier due to
    A[id] = big_calc4(id);                nowait
}←————— Implicit barrier
```

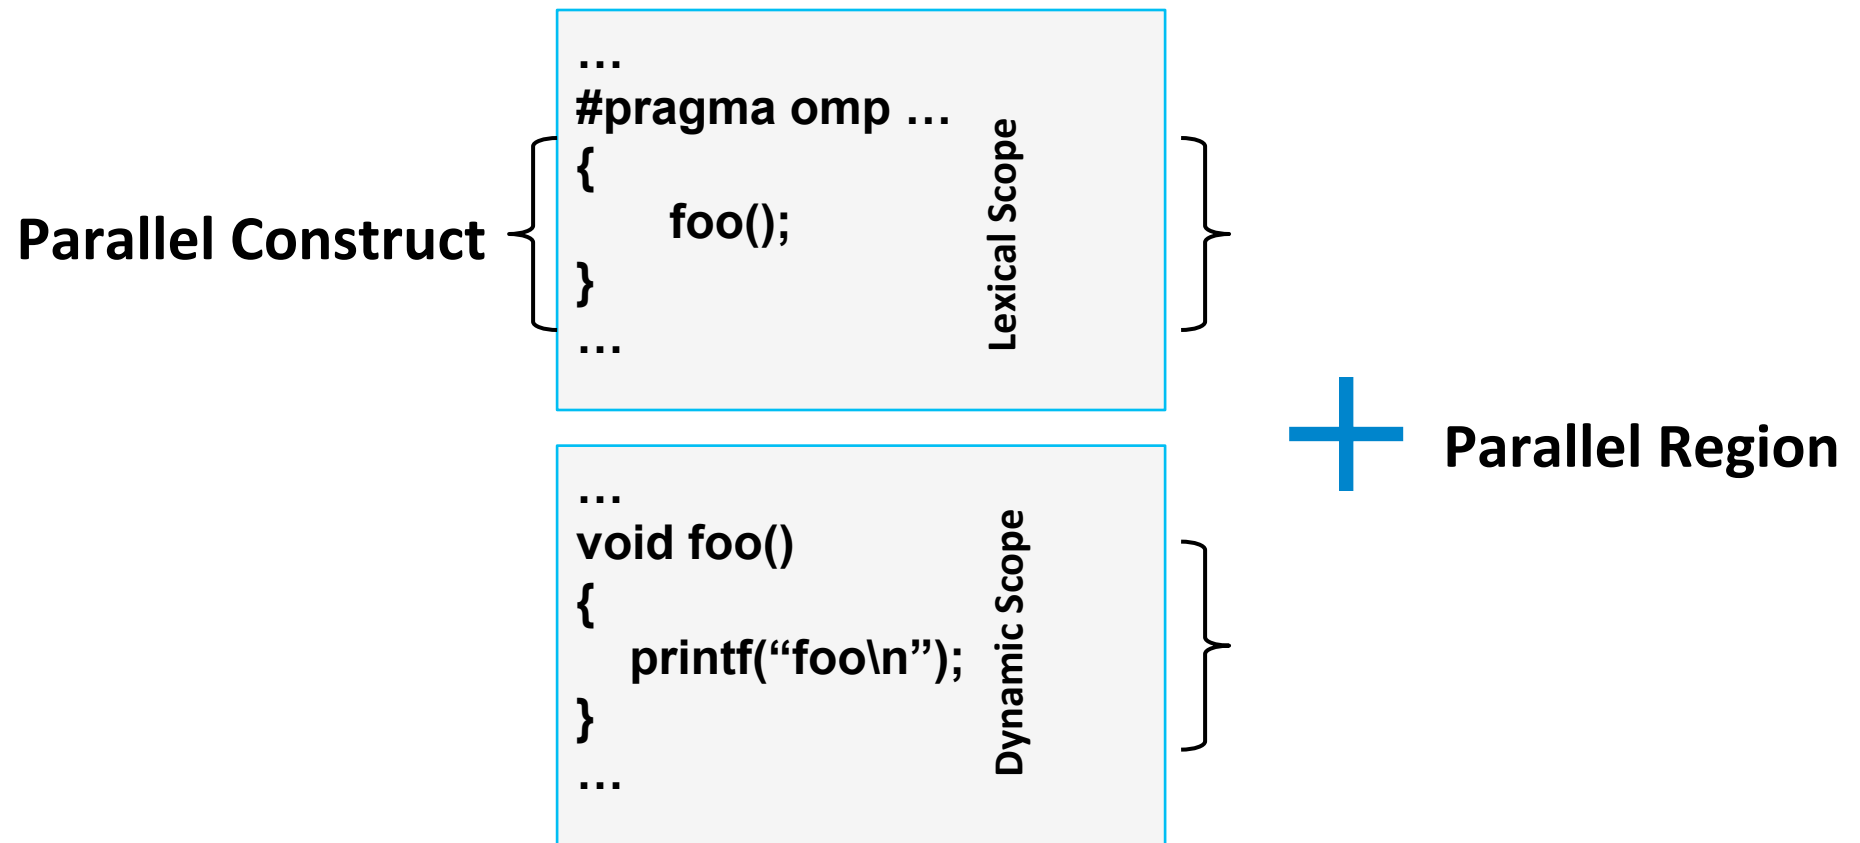
Barrier Example

- Barrier: Each thread waits until all threads arrive

```
!$omp parallel shared (A, B, C) private(id)

    id=omp_get_thread_num()
    A(id) = big_calc1(id)
    !$omp barrier
    !$omp do
    do i=0,N
        C(i)=big_calc3(i,A,N)
    end do ← Implicit barrier
    !$omp do
    do i=1,N
        B(i)=big_calc2(C, I, N)
    end do nowait ← No implicit barrier due to nowait
    A(id) = big_calc4(id)
!$omp end parallel ← Implicit barrier
```

Concepts



OpenMP Orphaned Work-sharing Constructs

- Work-sharing directives that appear outside the lexical extent of a parallel directive are called orphaned work sharing constructs.
(lexical extent: also called the “parallel construct”)
- When in the dynamic extent (i.e, within a called function or subroutine in a parallel region), the work-sharing construct behavior is identical (almost) to a work-sharing construct within the parallel region.
- When encountered from outside a parallel region (i.e. called from a serial portion of code) the master thread is the “team of threads”. It is safely invoked as serial code.

OpenMP Orphaned Work-sharing Constructs

```
!$omp parallel  
  call work(n,a,b,c)  
!$omp end parallel
```

```
  call work(n,a,b,c)
```

!-----

```
subroutine work(n,a,b,c)  
  use omp_lib  
  integer n, id, i  
  real*8, dimension(n):: a,b,c
```

```
  print*,omp_get_num_threads()
```

```
!$omp do  
  do i = 1,n  
    a(i) = b(i) + c(i)  
  end do
```

```
end subroutine
```

```
#pragma omp parallel  
  work(n,a,b,c);
```

```
  work(n,a,b,c);
```

//-----

```
int work(int n, double *a,...){
```

```
  int id, i;
```

```
  printf("%d\n",omp_get_num_threads());
```

```
  #pragma omp for  
  for(i=0; i<n; i++){  
    a[i] = b[i] + c[i];  
  }
```

```
}
```

OpenMP if clause

- Syntax of clause:
... **parallel if(*logical expression*)**
- Executes region as a parallel region if true, otherwise executed serially as a team of 1 thread.

```
!$omp      parallel if( N > 10000)  
#pragma omp parallel if( N > 10000)
```

modified

OpenMP critical region

- One thread at a time—
 - Can exist in parallel section, and in orphaned or serial code
 - Critical namespace is global. Same-named critical sections share a single lock.
- No guarantee for entry fairness, but
- Guaranteed forward processing
- Named critical region are independent **see below:

loop over i
↓
a(i) = worka(i)
add2b(a(i),b)
c(i) = workc(i)
add2d(c(i),d)

As separately named critical sections, B and D may have threads executing simultaneously.

B & D updates

```
!$omp parallel do  
do i = 1,n  
    a(i) = worka(i)  
    !$omp critical (B)  
    call add2b(a(i),b)  
    !$omp end critical (B)  
    c(i) = workc(i)  
    !$omp critical (D)  
    call add2d(c(i),d)  
    !$omp end critical (D)  
end do
```

```
#pragma omp parallel for  
{  
    for( i=0; i<n; i++)  
        a[i] = worka(i);  
    #pragma omp critical (B)  
        add2b(&a[i],&b);  
  
    c[i] = workc(i);  
    #pragma omp critical (D)  
        add2d(&c[i],&d);  
}
```

OpenMP flush directive

- Syntax: `!$omp flush [list()]`
`#pragma omp flush [list()]`
- A memory fence that inhibits movement of memory operations across the synchronization point.
- Point where executing thread has a consistent view of memory (of shared variables)
 - All memory operations (read/write) before synch. pt. must be performed before synch. pt. (no store later)
 - Likewise, all memory operations after synch. pt. must occur after synch. pt. (e.g. no prefetching across fence)
- Implicit flush (**all shared variables, no list of variables**)
 - Barrier, Entry/Exit Parallel, Critical & Ordered regions
 - Exit of worksharing region (except nowait), locks/tasks

OpenMP directives & clauses

Clause

Directive						
	parallel	do/for	sections	single	parallel do/for	parallel sections
if	X				X	X
default	X				X	X
shared	X				X	X
private	X	X	X	X	X	X
firstprivate	X	X	X	X	X	X
lastprivate		X	X		X	X
reduction	X	X	X		X	X
schedule		X			X	
nowait		X	X	X		

These don't accept clauses:

master
critical
barrier
atomic
flush

Table: Acceptable clauses for directives.

OpenMP runtime

Name	Type	Chunk	Chunk Size	Number of Chunks	Static or Dynamic	Computer Overhead
Simple Static	simple	no	N/P	P	static	lowest
Interleaved	simple	yes	C	N/C	static	low
Simple dynamic	dynamic	optional	C	N/C	dynamic	medium
Guided	guided	optional	decreasing from N/P	fewer than N/C	dynamic	high
Runtime	runtime	no	varies	varies	varies	varies

Guided (p =#of threads, n =# of iterations): starts at n/p , next thread get $\#unassigned/p$, etc. and decrements until chunk size is reached.