

Parallel Programming Using MPI

David Porter & Drew Gustafson

(612) 626-0802
help@msi.umn.edu

October 20, 2016

acroread /home/dhp/public/mpi.pdf

Agenda

- 10:00-10:15 Introduction to MSI Resources
- 10:15-10:30 Introduction to MPI
- 10:30-11:30 Blocking Communication
- 11:30-12:00 Hands-on
- 12:00- 1:00 Lunch
- 1:00- 1:45 Non-Blocking Communication
- 1:45- 2:20 Collective Communication
- 2:20- 2:45 Hands-on
- 2:45- 2:50 Break
- 2:50- 3:30 Collective Computation and Synchronization
- 3:30- 4:00 Hands-on

Introduction



Mesabi

HP Linux Cluster

750+ compute nodes

Each node has 2 x 12-core 2.5 GHz Intel Haswell processors

18,750+ cores

711+ Tflop aggregate performance

From 64 GB to 1 TB of memory per node

Aggregate memory: 67+ TB of RAM

40 GPU nodes:

2 Nvidia Tesla K40 GPUs

FDR/EDR Infiniband interconnect

➔ 5+ GB/s node-to-node communication

IB connect to Panasuas global file system



• <https://www.msi.umn.edu/content/mesabi>

• Supercomputing Institute

• for Advanced Computational Research



UNIVERSITY OF MINNESOTA

Driven to DiscoverSM

Itasca



HP Linux Cluster

1091 compute nodes

2 quad-core 2.8 GHz Intel Nehalem processors

24 GB of memory per node

Total of 8,728 cores

Aggregate of 26 TB of RAM

QDR Infiniband interconnect

→ 3+ GB/s node-to-node communication

• <https://www.msi.umn.edu/content/itasca>

Introduction to parallel programming

Serial → **one statement at a time**
one thread of execution, and
one process

Parallel → **multiple concurrent statements**
multiple threads of execution, and/or
one or more processes

Parallel Programming

Involves:

Decomposing work into many tasks
Distributing tasks to multiple threads or processes
Threads/processes work simultaneously
Coordinating work and communication of threads

Considerations

Type of parallel architecture being used
Type of communication needed between tasks

Parallel Programming

Uses

Multiple processors & threads

Multiple cores

Network (distributed memory machines, cluster, etc.)

Environment to create and manage parallel processing

Operating System

Parallel Programming Paradigms

Distributed memory: multiple processes

MPI

Shared Memory: multiple threads

OpenMP

Hardware Considerations

Memory architectures

Shared Memory (NUMA)

Distributed Memory

Cluster of Shared Memory "nodes"

Inter-node communication is required to:


Convey information and data between nodes

Start, stop, & synchronize processes across nodes

Shared Memory

Only one processor can access the shared memory location at a time.

Synchronization achieved by controlling tasks reading from and writing to the shared memory



Advantages:

Easy for user to use efficiently, data sharing among tasks is fast, ...

Disadvantages:


Memory is bandwidth limited, Total memory limited to one node

Distributed Memory

Data is shared across network using message passing

User code drives communication

Distributed Memory




Advantages:

Scalability, Each processor can rapidly access its own memory without interference

Disadvantages:

Programmer responsible for send/receive data between processes

Compute Cluster



Message Passing

MPI : Message Passing Interface

- A message passing library specification
- Model for distributed memory platforms
- Not a compiler
- For multi-core, clusters, and heterogeneous networks
- Permits development of parallel software libraries
- Provides access to advanced parallel hardware
- End uses

Applications
Libraries
Toolkits

MPI

- Widely accepted standard for distributed memory computing
- Support by all major vendors
- Efficient implementations exist for most parallel hardware
- Code that uses MPI is highly portable
- Very extensive and flexible interface that leaves most of the implementation details up to vendors
- Just a small subset of the functions (6 routines) can be used to write many applications

Parallel programming paradigms

SPMD (Single Program Multiple Data)

- All processes follow essentially the same execution path
- Data-driven execution

MPMD (Multiple Programs Multiple Data)

- Master and slave processes follow distinctly different execution paths
- Heterogeneous computing (GPU, PHI, ...)

MPI supports both

MPI Blocking Communication




Sending and Receiving Messages

Basic message passing :

One process send a message

Another process receives the message.



Questions:

- To whom is data sent?
- Where is the data?
- What type of data is sent?
- How much data is sent?
- How does the receiver identify it?

Message is divided into **data** and **envelope**

data

- buffer
- count
- data type

envelope

- process identifier (source/destination rank)
- message tag
- communicator

MPI Calling Conventions

Fortran Bindings:

Call **MPI_XXXX** (... , ierror)

- Case insensitive
- Almost all MPI calls are subroutines
- ierror is always the last parameter
- Program must include 'mpif.h'

C Bindings:

int ierror = **MPI_Xxxxxx** (...)

- Case sensitive (as it always is in C)
- All MPI calls are functions: most return integer error code
- Program must include "mpi.h"
- Parameters are passed by value → pass pointers to data buffers

MPI Basic Send/Receive

Blocking send:

MPI_Send (buffer, count, datatype, dest, tag, comm)

Blocking receive:

MPI_Recv (buffer, count, datatype, source, tag, comm, **status)**

MPI C Datatypes

MPI datatype	C datatype
MPI_CHAR	char
MPI_SHORT	signed short int
MPI_INT	signed int
MPI_LONG	signed long int
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED_SHORT	unsigned short int
MPI_UNSIGNED_LONG	unsigned long int
MPI_UNSIGNED	unsigned int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double
MPI_BYTE	byte
MPI_PACKED	

MPI Fortran Datatypes

MPI FORTRAN	FORTTRAN datatypes
MPI_INTEGER	INTEGER
MPI_REAL	REAL
MPI_REAL8	REAL*8
MPI_DOUBLE_PRECISION	DOUBLE PRECISION
MPI_COMPLEX	COMPLEX
MPI_LOGICAL	LOGICAL
MPI_CHARACTER	CHARACTER
MPI_BYTE	
MPI_PACKED	

MPI Process Identifier

- **MPI Application:** runs on a group of processes.
- **RANK:** one processes in this group
- **Rank NUMBER:** unique number for the process

In MPI communication:

- **Destination** is specified by **rank number**
- **Can point to all ranks:** `MPI_ANY_SOURCE`
- Processes are named according to their rank in the group
- Can have more than one group in an MPI application
- Groups are pointed to by a “communicator”

MPI Communicators

- A communicator
denotes a group of processes in an MPI application
- `MPI_COMM_WORLD`
predefined communicator
includes all processes in an MPI application
- New communicators
can be created in an MPI program
can point to some or all MPI “ranks”
can point to a re-ordering of ranks
- Most MPI programs only use `MPI_COMM_WORLD`

MPI Message Tag

Tags allow programmers to

- Organize / classify MPI messages
- Distinguish messages from the same source

The MPI standard guarantees that tags are

- integers in the range 0 ~ 32,767 (at least)
- most implementations allow a much larger range of tags
- upper bound on tag value: `MPI_TAG_UB`

`MPI_ANY_TAG` can be used as a wild card

MPI Blocking Communication Semantics

- MPI_SEND does not complete until buffer is empty (available for reuse)
- MPI_RECV does not complete until buffer is full (available for use)
- Completion of communication generally depends on the message size, system memory & network
- Blocking communication is simple to use but can be slow or cause deadlocks (if you are not careful).
- A blocking or nonblocking send can be paired to a blocking or nonblocking receive

Fortran Example

```
program MPI_small
include 'mpif.h'
integer rank, size, ierror, tag, status(MPI_STATUS_SIZE)
character(12) message

call MPI_INIT(ierror)
call MPI_COMM_SIZE(MPI_COMM_WORLD, size, ierror)
call MPI_COMM_RANK(MPI_COMM_WORLD, rank, ierror)
tag = 100
if(rank .eq. 0) then
    message = 'Hello, world'
    do i=1, size-1
        call MPI_SEND(message, 12, MPI_CHARACTER, i, tag,
                     MPI_COMM_WORLD, ierror)
    enddo
Else
    call MPI_RECV(message, 12, MPI_CHARACTER, 0, tag,
                 MPI_COMM_WORLD, status, ierror)
endif
print*, 'node' rank, ':', message
call MPI_FINALIZE(ierror)
end
```

C Example

```
#include <stdio.h>
#include "mpi.h"
main(int argc, char **argv)
{
    int rank, size, tag, rc, i;
    MPI_Status status;
    char message[20]

    rc = MPI_Init(&argc,&argv)
    rc = MPI_Comm_size(MPI_COMM_WORLD,&size);
    rc = MPI_Comm_rank(MPI_COMM_WORLD,&rank);
    tag = 100;
    if(rank == 0) {
        strcpy(message, "Hello, world");
        for (i=1; i<size; i++)
            rc = MPI_Send(message, 13, MPI_CHAR, i, tag, MPI_COMM_WORLD);
    }
    else
        rc = MPI_Recv(message, 13, MPI_CHAR, 0, tag, MPI_COMM_WORLD, &status);

    print("node %d : %.13s\n", rank,message);
    rc = MPI_Finalize();
}
```

Hands-on

<https://www.msi.umn.edu/content/mpi-hands-workshop>

Get example and build

```
cp -r /home/tech/public/examples/hello_mpi .  
cd hello_mpi  
module load intel impi  
make
```

Run interactively

```
mpirun -np 4 ./hello
```

Set the following for large-memory jobs

```
ulimit -s unlimited
```

MPI Non-Blocking Communication



Blocking Communication

`MPI_SEND` does not complete until buffer is copied out. Time depends on send mode and MPI implementation.

`MPI_RECV` does not complete until the message is completely received.

Completion of communication generally depends on the message size and the system buffer size.

Blocking communication is simple to use but may be prone to deadlocks.

Initiation of blocking communication may suffer from high latency → poor scaling to many MPI ranks.

Deadlocks

Two or more processes (MPI ranks) wait for each other to act before they act. They each stop before getting to the part of the code where they would have taken the action needed for other ranks to keep going.

Common example:

Two ranks call a blocking `mpi_recv` to each other.
Each waits for data from the other.
Neither ever sends it.

To avoid deadlocks

Different ordering of calls between ranks
Non-blocking calls
Use of `MPI_SendRecv`
Buffered mode

Send modes

Standard Mode (MPI_Send)

The standard MPI Send, the send will not complete until it is safe to modify the send buffer: buffer has at least been copied to a supplied system buffer. MPI may or may not buffer: depends on many details.

Synchronous mode (MPI_Ssend)

The send does not complete until after a matching receive has been posted

Buffered mode (MPI_Bsend)

User supplied buffer space is used for system buffering.

The send will complete as soon as the send buffer is copied to the user supplied buffer.

Ready mode (MPI_Rsend)

The send will send eagerly under the assumption that a matching receive has already been posted (an error results otherwise).

Non Blocking Communication

- **Non-blocking calls return immediately**
- **A completion call is needed to ensure the operation is finishes.**

For example:

```
MPI_ISEND( start, count, datatype, dest, tag, comm, request1)  
MPI_WAIT( request1, status )
```

```
MPI_IRECV( start, count, datatype, src, tag, comm, request2)  
MPI_WAIT( request2, status)
```

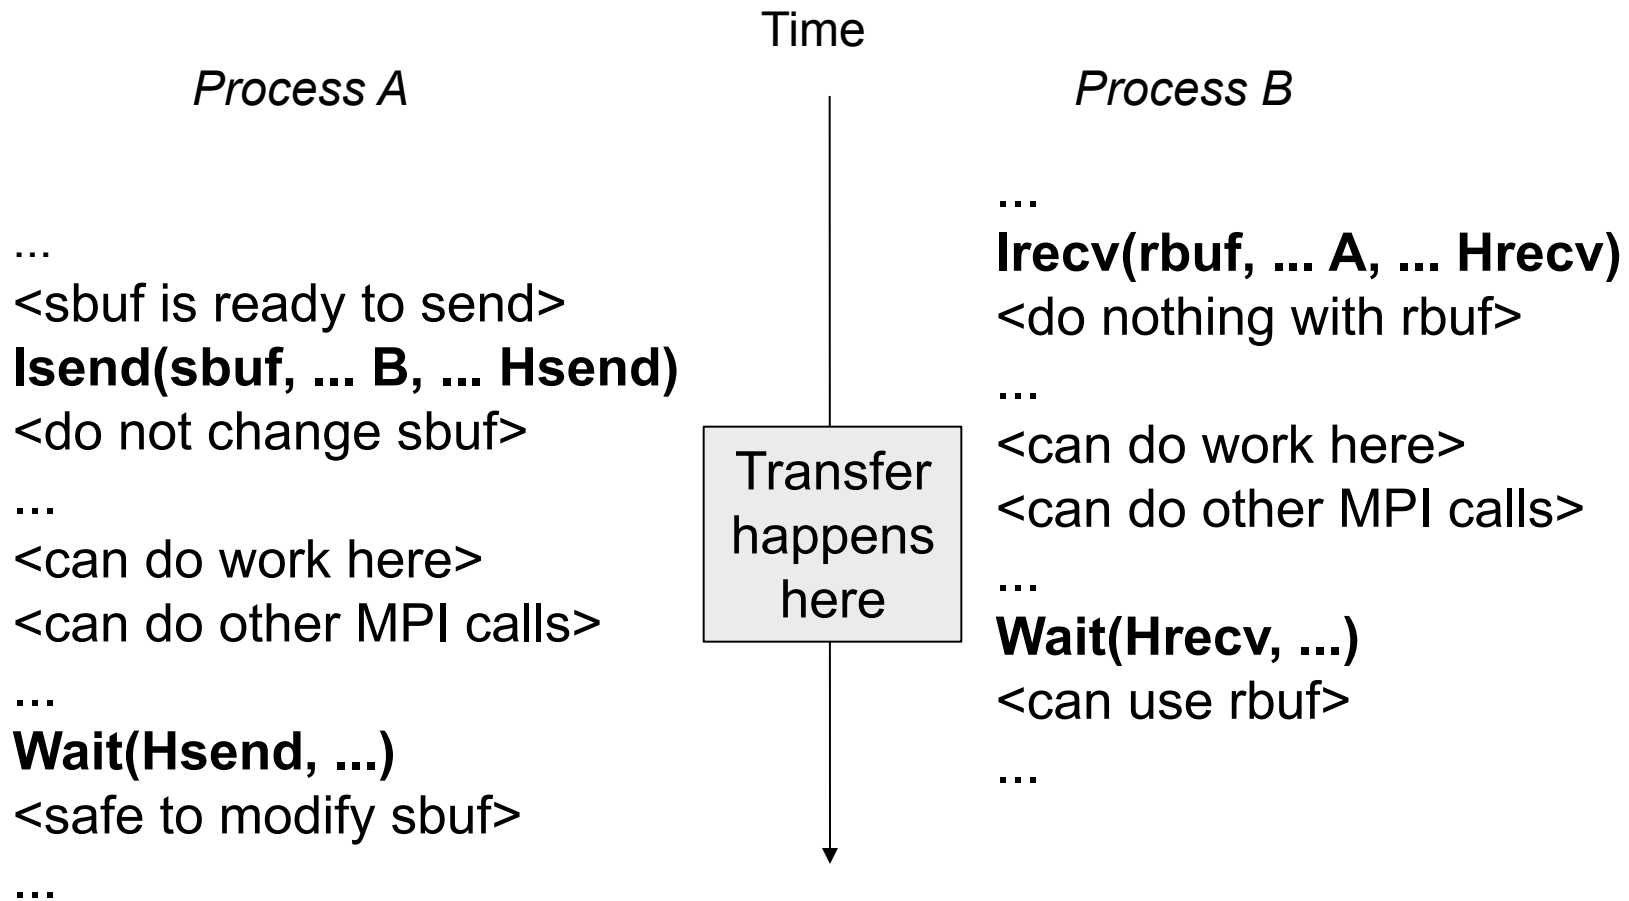
Or use for all non-blocking communications

```
MPI_WAITALL (count, request_array, status_arrsy)
```

One can also test the status without waiting using MPI_TEST

```
MPI_TEST(request, flag, status)
```

Non Blocking Communication Example



Illustrates conservative use of standard Isend & Irecv calls

Non-blocking Send Syntax

C:

```
int MPI_Isend(void* buf, int count, MPI_Datatype datatype, int dest, int tag,  
             MPI_Comm comm, MPI_Request *request)
```

FORTTRAN:

```
MPI_ISEND(BUF, COUNT, DATATYPE, DEST, TAG,  
          COMM, REQUEST, IERROR)
```

[IN buf] initial address of send buffer (any data type)
[IN count] number of elements in send buffer (integer)
[IN datatype] datatype of each send buffer element (defined constant)
[IN dest] rank of destination (integer)
[IN tag] message tag (integer)
[IN comm] communicator (handle)
[OUT request] communication request (handle)

Non-blocking Recv Syntax

C:

```
int MPI_Irecv(void* buf, int count, MPI_Datatype datatype, int source, int tag,  
             MPI_Comm comm, MPI_Request *request)
```

FORTRAN:

```
MPI_IRecv(BUF, COUNT, DATATYPE, SOURCE, TAG,  
          COMM, REQUEST, IERROR)
```

[OUT buf] initial address of receive buffer (choice)
[IN count] number of elements in receive buffer (integer)
[IN datatype] datatype of each receive buffer element (defined constant)
[IN dest] rank of source (integer)
[IN tag] message tag (integer)
[IN comm] communicator (handle)
[OUT request] communication request (handle)

Non-blocking Communication completion calls

Wait: **MPI_WAIT** or **MPI_WAITALL**

Used for non-blocking Sends and Receives
Suspends until an operation completes

MPI_WAIT syntax

Fortran call MPI_WAIT (request, status, ierror)
C: ierror = MPI_Wait (request, status)

Test: **MPI_TEST**

Returns immediately with information about a non-blocking send or receive. Gives immediate answer to: **is send or receive done?**

MPI_TEST Syntax

Fortran: call MPI_TEST (request, flag, status, ierror)
C: ierror = MPI_Test (request, flag, status)

Non-blocking Communication completion calls

A request object can be deallocated at any time

Use the following operation:

```
MPI_REQUEST_FREE(request)  
[ INOUT request ] communication request (handle)
```

```
C:      ierror = MPI_Request_free(MPI_Request *request)
```

```
FORTRAN:      call MPI_REQUEST_FREE(REQUEST, IERROR)
```

Non-blocking Communication Examples

Example: Simple usage of nonblocking operations and MPI_WAIT

```
IF(rank.EQ.0) THEN
    CALL MPI_ISEND(a(1), 10, MPI_REAL, 1, tag, comm, request, ierr)
    ****do some computation to mask latency****
    CALL MPI_WAIT(request, status, ierr)
ELSE
    CALL MPI_Irecv(a(1), 10, MPI_REAL, 0, tag, comm, Request, ierr)
    ****do some computation to mask latency****
    CALL MPI_WAIT(Request, status, ierr)
END IF
```

```
INCLUDE "mpif.h"  
INTEGER ierror, rank, size, status(MPI_STATUS_SIZE), requests (2)
```

```
CALL MPI_INIT(ierror)  
CALL MPI_COMM_SIZE(MPI_COMM_WORLD, np, ierr)  
CALL MPI_COMM_RANK(MPI_COMM_WORLD, rank, ierr)  
IF(rank.eq.0) THEN  
  c = 9.2  
  a = 4.2  
  b = 8.4  
  CALL MPI_ISEND(a, 1, MPI_REAL,1,101,MPI_COMM_WORLD, requests(1), ierror)  
  ! Can do computations which do not overwrite a  
  b = b + a  
  CALL MPI_WAIT(requests(1), status, ierror)  
  d = b + c  
ELSE  
  a = 14.2  
  s = 18.4  
  CALL MPI_IRECV(c,1,MPI_REAL,0,101,MPI_COMM_WORLD, requests(2), ierror)  
  ! Do not read from or overwrite c till wait  
  CALL MPI_WAIT(requests(2), status,ierror)  
  c = a + c  
END IF  
  
CALL MPI_FINALIZE(ierror)  
STOP  
END
```


Non-blocking Communication

Gain

- Avoid Deadlocks
- Decrease Synchronization Overhead
- Can Reduce System Overhead
- Post non-blocking sends/receives early and do waits late
- Recommended: do MPI_Irecv before the MPI_Rsend is called.

Be careful with reads and writes

- Avoid writing to send buffer between MPI_Isend and MPI_Wait
- Avoid reading or writing in receive buffer between MPI_Irecv and MPI_Wait

MPI

Collective Communication



MPI Collective Communication

Three Classes of Collective Operations

Data movement:

broadcast, gather, all-gather, scatter, and all-to-all

Collective computations (reductions):

Data from all members in a group is “reduced” to produce a global result (min, max, sum, ...).


Synchronization:

processes wait until all members of the group have reached the synchronization point

Every process must call the same collective communication function.

MPI Collective Communication Broadcast

One rank sends (broadcasts) a block of data to all the ranks in a group.



MPI Collective Communication Broadcast

Syntax

C: `int MPI_Bcast (void* buffer, int count, MPI_Datatype datatype, int root, MPI_Comm comm)`

Fortran: `MPI_BCAST (buffer, count, datatype, root, comm, ierr)`

where:

buffer: is the starting address of a buffer

count: is an integer indicating the number of data elements in the buffer

datatype: is MPI defined constant indicating the data type of the elements

in the buffer

root: is an integer indicating the rank of broadcast root process


comm: is the communicator

The MPI_BCAST must be called by each process in the group, specifying the same comm and root. The message is sent from the root process to all processes in the group.

MPI Collective Communication

Gather

Data is distributed throughout all processors in the group.
Collect distributed data to a specified process (rank).



```
real a(100), rbuf(MAX)
```

```
call mpi_gather(a, 100, MPI_REAL, rbuf, 100, MPI_REAL, root, comm, ierr)
```

MPI Collective Communication

Gather

Syntax

C:

```
int MPI_Gather(void* sbuf, int scount, MPI_Datatype stype, void* rbuf, int rcount,  
              MPI_Datatype rtype, int root, MPI_Comm comm)
```

FORTTRAN:

```
MPI_GATHER (sbuf, scount, stype, rbuf, rcount, rtype, root, comm, ierr)
```

where:

sbuf:	is the starting address of a buffer,
scount:	is the number of elements in the send buffer,
stype:	is the data type of send buffer elements,
rbuf:	is the address of the receive buffer
rcount:	is the number of elements for any single receive
rtype:	is the data type of the receive buffer elements
root:	is the rank of receiving process, and
comm:	is the communicator
ierr:	is error message

MPI Collective Communication

Gather Example

```

INCLUDE 'mpif.h'
DIMENSION A(25, 100), b(100), cpart(25), ctotat(100)
INTEGER root, rank
CALL MPI_INIT(ierr)
CALL MPI_COMM_SIZE(MPI_COMM_WORLD, np, ierr)
CALL MPI_COMM_RANK(MPI_COMM_WORLD, rank, ierr)
root=1
Call set_a(A,rank)
Call set_b(B)

DO I=1,25
  cpart(I)=0.
  DO K=1,100
    cpart(I)=cpart(I)+A(I,K)*b(K)
  END DO
END DO

CALL MPI_GATHER (cpart, 25, MPI_REAL, ctotat, 25,
&               MPI_REAL, root, MPI_COMM_WORLD, ierr)
If(rank.eq.root) print*, (ctotat(I),I=1,100)
CALL MPI_FINALIZE(ierr)
END

```


A	*	b	=	c
Process		1		1
Process		2		2
Process		3		3
Process		4		4

- A: Matrix distributed by rows
- B: Vector shared by all process
- C: results to get by the root process

MPI Collective Communication

Scatter

Distribute data of to all the processes (ranks) in a group.



MPI Collective Communication

Scatter Syntax

C:

```
int MPI_Scatter(void* sbuf, int scount, MPI_Datatype stype, void* rbuf,  
               int rcount, MPI_Datatype rtype, int root, MPI_Comm comm)
```

FORTRAN:

```
MPI_SCATTER(sbuf, scount, stype, rbuf, rcount, rtype, root, comm, ierr)
```

where:

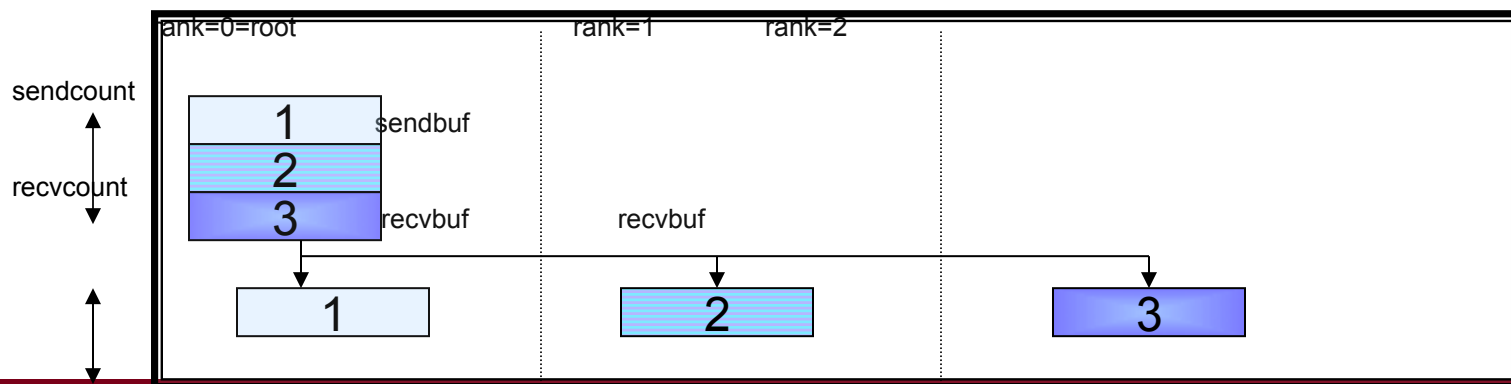
sbuf:	is the address of the send buffer,
scount:	is the number of elements sent to each process,
stype:	is the data type of the send buffer elements,
rbuf:	is the address of the receive buffer,
rcount:	is the number of elements in the receive buffer,
rtype:	is the data type of the receive buffer elements,
root:	is the rank of the sending process, and
comm:	is the communicator

Note: sbuf is significant for root process only

Sample execution

```
PROGRAM scatter
INCLUDE 'mpif.h'
INTEGER isend(3)
CALL MPI_INIT(ierr)
CALL MPI_COMM_SIZE(MPI_COMM_WORLD, nprocs, ierr)
CALL MPI_COMM_RANK(MPI_COMM_WORLD, myrank, ierr)
IF (myrank == 0) THEN
    DO i=1, nprocs
        isend (i) = i
    ENDDO
    end if
CALL MPI_SCATTER (isend, 1, MPI_INTEGER,
&                irecv, 1, MPI_INTEGER, 0,
&                MPI_COMM_WORLD, ierr)
PRINT *, 'irecv = ', irecv
CALL MPI_FINALIZE(ierr)
END
```

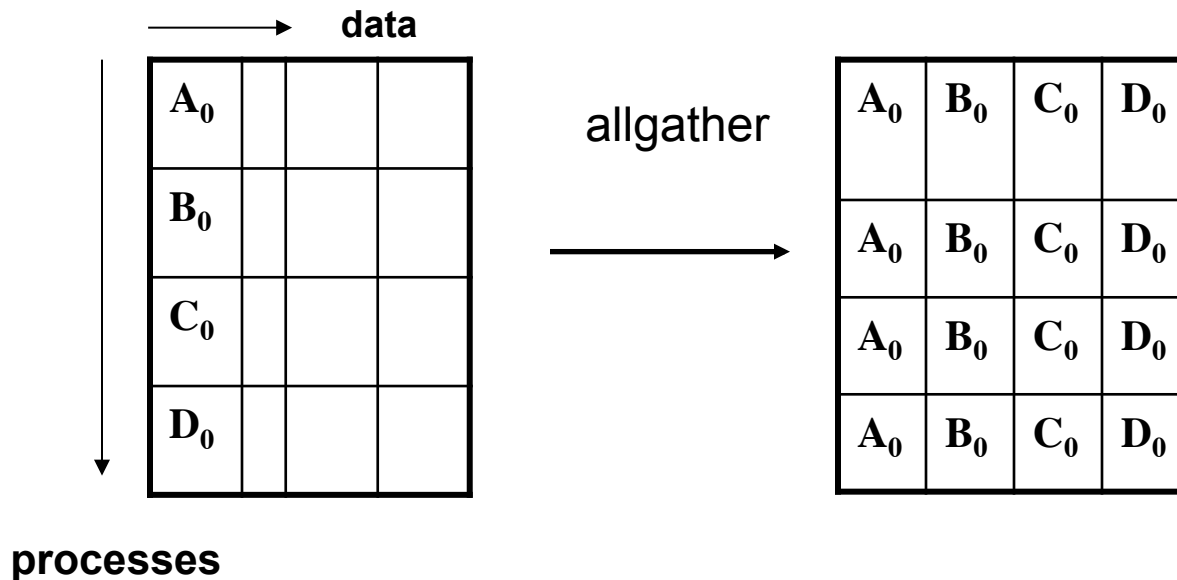
```
$ mpirun -np 3 ./a.out
0: irecv = 1
1: irecv = 2
2: irecv = 3
```



MPI Collective Communication

All Gather

MPI_ALLGATHER can be thought of as MPI_GATHER where all processes, not only one, receive the result.



The syntax of MPI_ALLGATHER is similar to MPI_GATHER. However, the argument root is dropped

MPI Collective Communication

All Gather Syntax

C:

```
int MPI_Allgather (void* sbuf, int scout, MPI_Datatype stype, void* rbuf, int rcount,
                  MPI_Datatype rtype, MPI_Comm comm)
```

FORTRAN

```
MPI_ALLGATHER (sbuf, scout, stype, rbuf, rcount, rtype, comm, ierr)
```

Example: back to the previous “gather” example, what should we do if every process needs the results of array Ctotal for next computation?

Replace

```
CALL MPI_GATHER (cpart, 25, MPI_REAL, ctotal, 25,
                MPI_REAL, root, MPI_COMM_WORLD, ierr)
```

With

```
CALL MPI_ALLGATHER (cpart, 25, MPI_REAL, ctotal, 25,
                   MPI_REAL, MPI_COMM_WORLD, ierr)
```

MPI Collective Communication

Alltoall

Send Buffer

Data →

**P
r
o
c
e
s
s
e
s**

A ₀	B ₀	C ₀	D ₀	E ₀	F ₀
A ₁	B ₁	C ₁	D ₁	E ₁	F ₁
A ₂	B ₂	C ₂	D ₂	E ₂	F ₃
A ₃	B ₃	C ₃	D ₃	E ₃	F ₃
A ₄	B ₄	C ₄	D ₄	E ₄	F ₄
A ₅	B ₅	C ₅	D ₅	E ₅	F ₅

Receive Buffer

Data →

**P
r
o
c
e
s
s
e
s**

A ₀	A ₁	A ₂	A ₃	A ₄	A ₅
B ₀	B ₁	B ₂	B ₃	B ₄	B ₅
C ₀	C ₁	C ₂	C ₃	C ₄	C ₅
D ₀	D ₁	D ₂	D ₃	D ₄	D ₅
E ₀	E ₁	E ₂	E ₃	E ₄	E ₅
F ₀	F ₁	F ₂	F ₃	F ₄	F ₅

MPI Collective Communication

Alltoall Syntax

C:

```
int MPI_Alltoall(void* sbuf, int scount, MPI_Datatype stype, void* rbuf, int  
                rcount, MPI_Datatype rtype, MPI_Comm comm )
```

FORTRAN:

```
MPI_ALLTOALL (sbuf, scount, stype, rbuf, rcount, rtype, comm, ierr)
```

where:

sbuf: is the starting address of send buffer,
scount: is the number of elements sent to each process,
stype: is the data type of send buffer elements,
rbuf: is the address of receive buffer,
rcount: is the number of elements received from any process,
rtype: is the data type of receive buffer elements, and
comm: is the group communicator.




Figure of MPI_ALLTOALL

```
$ a.out -procs 3
0: isend 1 2 3
1: isend 4 5 6
2: isend 7 8 9
0: irecv 1 4 7
1: irecv 2 5 8
2: irecv 3 6 9
```

```
PROGRAM alltoall
INCLUDE 'mpif.h'
INTEGER isend (3), irecv (3)
CALL MPI_INIT(ierr)
CALL MPI_COMM_SIZE(MPI_COMM_WORLD, nprocs, ierr)
CALL MPI_COMM_RANK(MPI_COMM_WORLD,
& myrank, ierr)
DO i=1, nprocs
    isend (i) = i + nprocs * myrank
ENDDO
PRINT *, 'isend =', isend
CALL MPI_ALLTOALL(isend, 1, MPI_INTEGER,
&    irecv, 1, MPI_INTEGER,
&    MPI_COMM_WORLD, ierr)
PRINT *, 'irecv =', irecv
CALL MPI_FINALIZE(ierr)
END
```


Hands-on

<https://www.msi.umn.edu/content/mpi-hands-workshop>

Get example and build

```
cp -r /home/tech/public/examples/hello_mpi .  
cd hello_mpi  
module load intel impi  
make
```

Run interactively

```
mpirun -np 4 ./hello
```

Set the following for large-memory jobs

```
ulimit -s unlimited
```

MPI

Collective Computations and Synchronization



MPI_Reduce

These routines perform a global operation across all members of a group

The partial result in each process in the group is combined in one specified process or all the processes using some desired function.

Three reduces routines:

MPI_REDUCE returns results to a single process;

MPI_ALLREDUCE returns results to all processes in the group;

MPI_REDUCE_SCATTER scatters a vector, which results in a reduce operation, across all processes.

Fortran

MPI_REDUCE (sbuf, rbuf, count, stype, op, root, comm, ierr)

MPI_ALLREDUCE (sbuf, rbuf, count, stype, op, comm, ierr)

MPI_REDUCE_SCATTER (sbuf, rbuf, rcounts, stype, op, comm, ierr)

where

sbuf: is the address of send buffer,
rbuf: is the address of receive buffer,
count: is the number of elements in send buffer,
stype: is the data type of elements of send buffer,
op: is the reduce operation (which may be MPI predefined, or your own),
root: is the rank of the root processes, and
comm: is the group communicator.

C:

int **MPI_Reduce** (void* sbuf, void* rbuf, int count, MPI_Datatype stype, MPI_Op op, int root, MPI_Comm comm)

int **MPI_Allreduce**(void* sbuf, void* rbuf, int count, MPI_Datatype stype, MPI_Op op, MPI_Comm comm)

int **MPI_Reduce_scatter** (void* sbuf, void* rbuf, int* rcounts, MPI_Datatype stype, MPI_Op op, MPI_Comm comm)

where

sbuf: is the address of send buffer,
rbuf: is the address of receive buffer,
count: is the number of elements in send buffer,
stype: is the data type of elements of send buffer,
op: is the reduce operation (which may be MPI predefined, or your own),
root: is the rank of the root processes, and
comm: is the group communicator.

MPI Predefined Reduce Operations

	Name	Meaning	C type	FORTTRAN type
MPI_MAX		maximum	integer, float	integer, real, complex
MPI_MIN		minimum	integer, float	integer, real, complex
MPI_SUM		sum	integer, float	integer, real, complex
MPI_PROD		product	integer, float	integer, real, complex
MPI LAND		logical and	integer	logical
MPI_BAND		bit-wise and	integer, MPI_BYTE	integer, MPI_BYTE
MPI_LOR		logical or	integer	logical
MPI BOR		bit-wise or	integer, MPI_BYTE	integer, MPI_BYTE
MPI_LXOR		logical xor	integer	logical
MPI_BXOR		bit-wise xor	integer MPI_BYTE	integer, MPI_BYTE
MPI_MAXLOC		max value and location	combination of int, float, double, and long double	combination of integer, real, complex, double precision
MPI_MINLOC		min value and location	combination of int, float, double, and long double	combination of integer, real complex, double precision

MPI_REDUCE

Usage: CALL **MPI_REDUCE** (sendbuf, recvbuf, count, datatype, op, root, comm, ierror)

Parameters

(CHOICE) **sendbuf** The address of the send buffer (IN)
(CHOICE) **recvbuf** The address of the receive buffer, sendbuf and recvbuf cannot overlap in memory. (significant only at root) (OUT)
INTEGER **count** The number of elements in the send buffer (IN)
INTEGER **datatype** The data type of elements of the send buffer (handle) (IN)
INTEGER **op** The reduction operation (handle) (IN)
INTEGER **root** The rank of the root process (IN)
INTEGER **comm** The communicator (handle) (IN)
INTEGER **ierror** The Fortran return code

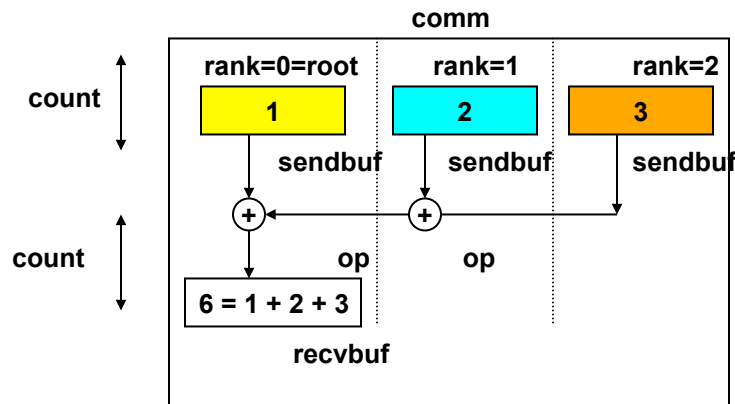


Figure: MPI_REDUCE for Scalar Variables

Sample Program

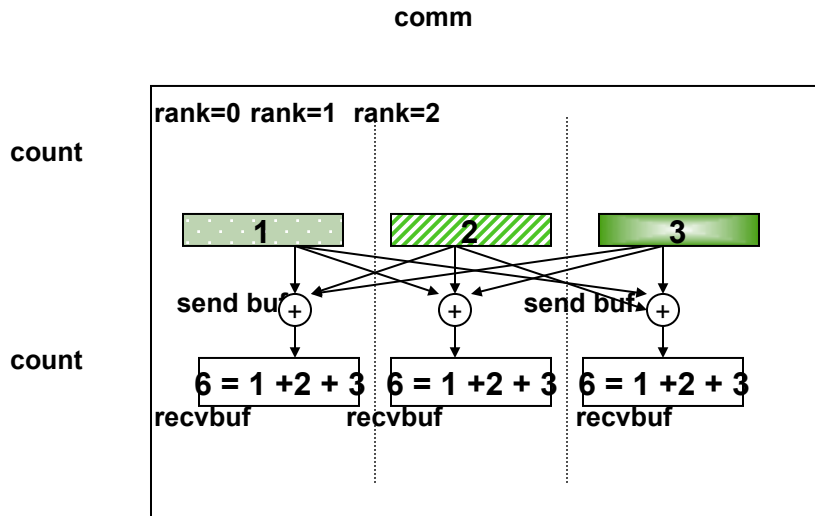
```
PROGRAM reduce
  include 'mpif.h'
  CALL MPI_INIT (ierr)
  CALL MPI_COMM_SIZE (MPI_COMM_WORLD, nprocs, ierr)
  CALL MPI_COMM_RANK (MPI_COMM_WORLD, myrank, ierr)
  isend= myrank+1
  CALL MPI_REDUCE (isend, irecv, 1, MPI_INTEGER,
&                MPI_SUM, 0, MPI_COMM_WORLD, ierr)
  IF (myrank==0) THEN
    PRINT *, 'irecv =', irecv
  endif
  CALL MPI_FINALIZE (ierr)
END
```

Sample execution

```
% a.out -procs 3
% 0: irecv = 6
```

MPI_ALLREDUCE

Usage: CALL **MPI_ALLREDUCE** (sendbuf, recvbuf, count, datatype, op, comm, ierror)



Sample program

```
PROGRAM allreduce
  INCLUDE 'mpif.h'
  CALL MPI_INIT (ierr)
  CALL MPI_COMM_SIZE(MPI_COMM_WORLD, nprocs, ierr)
  CALL MPI_COMM_RANK (MPI_COMM_WORLD, myrank, ierr)
  isend = myrank + 1
  CALL MPI_ALLREDUCE (isend, irecv, 1, MPI_INTEGER, MPI_SUM,
    & MPI_COMM_WORLD, ierr)
  PRINT *, 'irecv =', irecv
  CALL MPI_FINALIZE (ierr)
END
```

Parameters

(CHOICE) **sendbuf** The starting address of the send buffer (IN)

(CHOICE) **recvbuf** The starting address of the receive buffer,,
sendbuf and recvbuf cannot overlap in memory (OUT)

INTEGER **count** The number of elements in the send buffer (IN)

INTEGER **datatype** The data type of elements of the send buffer (handle)
(IN)

INTEGER **op** The reduction operation (handle)(IN)

INTEGER **comm** The communicator (handle) (IN)

INTEGER **ierror** The Fortran return code

Sample execution

\$ a.out -procs 3

0: irecv = 6
1: irecv = 6
2: irecv = 6

MPI_REDUCE_SCATTER

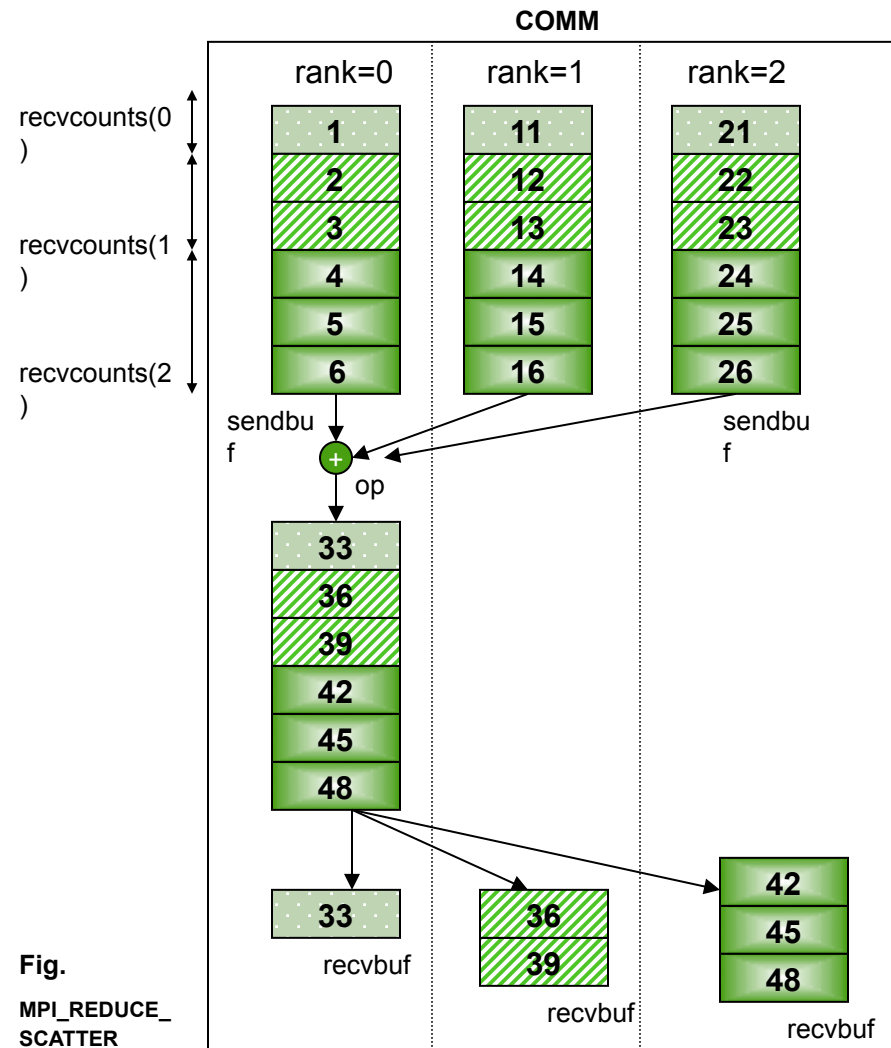
Usage: CALL MPI_REDUCE_SCATTER(sendbuf, recvbuf, recvcunts, datatype, op, comm, ierror)

Parameters

- (CHOICE) **sendbuf** The starting address of the send buffer (IN)
- (CHOICE) **recvbuf** The starting address of the receive buffer, sendbuf and recvbuf cannot overlap in memory. (OUT)
- INTEGER **recvcunts**(*)
Integer array specifying the number of elements in result distributed to each process. Must be identical on all calling processes. (IN)
- INTEGER **datatype** The data type of elements of the input buffer (handle) (IN)
- INTEGER **op** The reduction operation (handle) (IN)
- INTEGER **comm** The communicator (handle) (IN)
- INTEGER **ierror** The Fortran return code

Description MPI_REDUCE_SCATTER first performs an element-wise reduction on vector of count = $\sum_i \text{recvcunts}(i)$ elements in the send buffer defined by sendbuf, count and datatype. Next, the resulting vector is split into n disjoint segments, where n is the number of members in the group. Segment i contains recvcunts(i) elements. the ith segment is sent to process i and stored in the receive buffer defined by recvbuf, recvcunts(i) and datatype. MPI_REDUCE_SCATTER is functionally equivalent to MPI_REDUCE with count equal to the sum of recvcunts(i) followed by MPI_SCATTERV with sendcounts equal to recvcunts. All processes in comm need to call this routine.

CALL **MPI_REDUCE_SCATTER** (sendbuf, recvbuf, recvcounts, datatype, op, comm, ierror)



Sample Program

```

PROGRAM reduce_scatter
INCLUDE 'mpif.h'
INTEGER isend (6), irecv (3)
INTEGER ircnt (0:2)
DATA ircnt/1,2,3/
CALL MPI_INIT (ierr)
CALL MPI_COMM_SIZE(MPI_COMM_WORLD, nprocs, ierr)
CALL MPI_COMM_RANK(MPI_COMM_WORLD, myrank, ierr)
DO i=1,6
    isend (i) = i + myrank * 10
ENDDO
CALL MPI_REDUCE_SCATTER(isend, irecv, ircnt, MPI_INTEGER,
&                        MPI_SUM, MPI_COMM_WORLD,
ierr)
PRINT *, 'irecv =', irecv
CALL MPI_FINALIZE(ierr)
END
    
```

\$ a.out -procs 3

```

0: irecv = 33   0   0
1: irecv = 36  39   0
2: irecv = 42  45  48
    
```

Scan

A scan or prefix-reduction operation performs partial reductions on distributed data.

C: `int MPI_Scan (void* sbuf, void* rbuf, int count, MPI_Datatype datatype, MPI_Op op, MPI_Comm comm`

FORTTRAN: `MPI_SCAN (sbuf, rbuf, count, datatype, op, comm, ierr)`

Where:

sbuf: is the starting address of the send buffer,

rbuf: is the starting address of receive buffer,

count: is the number of elements in input buffer,

datatype: is the data type of elements of input buffer

op: is the operation, and

comm: is the group communicator.

Usage: CALL **MPI_SCAN** (sendbuf, recvbuf, count, datatype, op, comm, ierror)

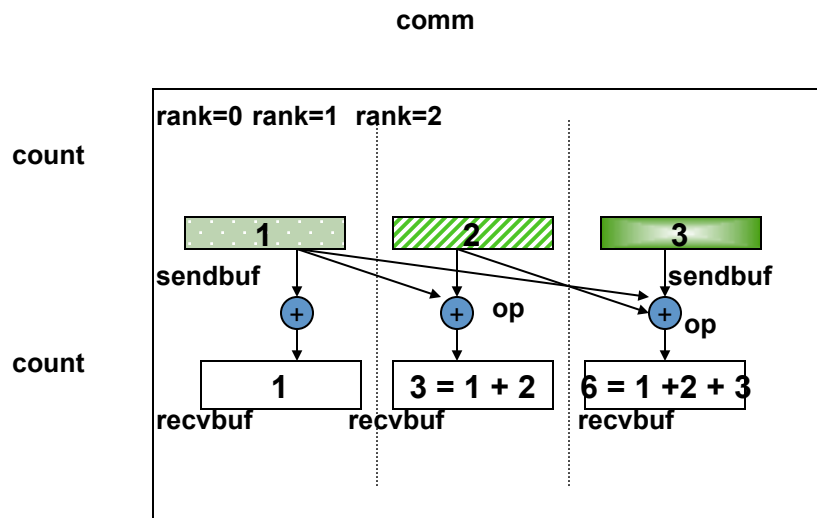


Figure. MPI_SCAN

Parameters

(CHOICE) sendbuf	The starting address of the send buffer (IN)
(CHOICE) recvbuf	The starting address of the receive buffer, sendbuf and recvbuf cannot overlap in memory (OUT)
INTEGER count	The number of elements in sendbuf (IN)
INTEGER datatype	The data type of elements of sendbuf (handle) (IN)
INTEGER op	The reduction operation (handle) (IN)
INTEGER comm	The communicator (handle) (IN)
INTEGER ierror	The Fortran return code

Sample program

```

PROGRAM scan
  INCLUDE 'mpif.h'
  CALL MPI_INIT (ierr)
  CALL MPI_COMM_SIZE (MPI_COMM_WORLD, nprocs, ierr)
  CALL MPI_COMM_RANK (MPI_COMM_WORLD, myrank, ierr)
  isend = myrank + 1
  CALL MPI_SCAN (isend, irecv, 1, MPI_INTEGER,
    & MPI_SUM, MPI_COMM_WORLD, ierr)
  PRINT *, 'irecv =' irecv
  CALL MPI_FINALIZE(ierr)
END

```

Sample execution

```

$ a.out -procs 3
0: irecv = 1
0: irecv = 3
0: irecv = 6

```

User-defined Operations

User can define his/her own reduce operation
Makes use of the `MPI_OP_CREATE` function

Performance Issues

A great deal of hidden communication takes place with collective communication. Performance depends greatly on the particular implementation of MPI. Because there may be forced synchronization, not always best to use collective communication.

Barrier Synchronization

Two types of synchronization:

Implicit synchronization

Explicit synchronization: `MPI_BARRIER`

MPI provides a function call, `MPI_BARRIER`, to synchronize all processes within a communicator.

A barrier is simply a synchronization primitive. A node calling it will be blocked until all the nodes within the group have called it.

Barrier Synchronization

The syntax of MPI_BARRIER for both C and Fortran program is:

- C:

MPI_Barrier (MPI_Comm comm)

- FORTRAN

MPI_BARRIER (comm, ierr)

where:

MPI_Comm:	is an MPI predefined structure of communicators,
comm:	is an integer denoting a communicator
ierr:	is an integer return error code.

Hands-on

<https://www.msi.umn.edu/content/mpi-hands-workshop>

Get example and build

```
cp -r /home/tech/public/examples/hello_mpi .  
cd hello_mpi  
module load intel impi  
make
```

Run interactively

```
mpirun -np 4 ./hello
```

Set the following for large-memory jobs

```
ulimit -s unlimited
```


THANK YOU

More info at
www.msi.umn.edu
612-626-0802

•The University of Minnesota is an equal opportunity educator and employer. This PowerPoint is available in alternative formats upon request. Direct requests to Minnesota Supercomputing Institute, 599 Walter library, 117 Pleasant St. SE, Minneapolis, Minnesota, 55455, 612-624-0528.

•© 2009 Regents of the University of Minnesota. All rights reserved.

•Supercomputing Institute
•for Advanced Computational Research



UNIVERSITY OF MINNESOTA
Driven to DiscoverSM